

The Physical and Contextual Limits of Quantum Speedup

Karl Svozil ¹, *

¹*Institute for Theoretical Physics, TU Wien, Wiedner Hauptstrasse 8-10/136, 1040 Vienna, Austria*
(Dated: May 12, 2026)

Quantum computation is often popularly described as the simultaneous execution of exponentially many classical computations. This article offers a conceptual clarification of why that “branchwise parallelism” picture is misleading: by branchwise parallelism we mean the idea that the components of a superposition behave as independently inspectable classical computational branches. Quantum speedups arise instead from reversible embeddings of algebraic structure made accessible through engineered interference patterns. We review this mechanism through several constraints: unitary garbage erasure is impossible, copying and deletion are context-dependent, and contextuality obstructs a single global classical history. We also distinguish circuit or unitary universality from Turing universality: dense generation of unitaries is not the same as symbolic computation over unbounded inputs with recursion, uniformity, and self-reference. In closed unitary dynamics there is no nontrivial absorbing halting state of the classical many-to-one kind; operational termination requires clocks, flags, measurements, open-system records, or external control. Exponential Hilbert-space dimension supplies a geometry for interference and high-dimensional embedding, not unlimited classical readout.

I. INTRODUCTION

A common description of quantum computation says that a quantum computer evaluates a function on many inputs “at once.” The image is seductive: a register is placed in a superposition, a function is applied to all components, and a final measurement extracts the desired answer. This picture is often invoked when discussing Deutsch–Jozsa [1], Simon’s algorithm [2], Shor’s algorithm [3], Grover search [4], and related oracle or Fourier-sampling procedures [5, 6].

Taken literally, however, this picture is wrong. A quantum computer does not return the list of all computed values. The components of a superposition cannot be individually inspected, copied, deleted, or archived as classical alternatives. They are amplitudes in a Hilbert-space vector, and only suitably arranged interference effects become accessible in a final measurement context, where a context is understood as a compatible set of observables, or equivalently a choice of measurement basis in the nondegenerate finite-dimensional case.

The thesis of this article is the following: Quantum speedups are not many classical computations performed in parallel. They arise when the relevant predicate of a task can be embedded into a unitary evolution so that it becomes accessible by interference in a chosen measurement context. At the same time, copying and erasure are physically constrained by no-cloning, no-deletion, and Landauer irreversibility, while contextuality prevents the intermediate process from being interpreted as a single global classical history or as a noncontextual value assignment.

A related ambiguity concerns the word “universal.” Circuit or unitary universality means that a finite gate set

can approximate a large class of unitary transformations. Turing universality, by contrast, concerns the symbolic simulation of arbitrary effective procedures on inputs of unbounded length. These are not the same notion. A family of quantum circuits may define an algorithm in the complexity-theoretic sense only when a uniform classical or quantum rule generates the appropriate circuit for each input size. A single finite circuit, or even a universal finite gate set, is not by itself equivalent to the class of partial recursive functions.

This formulation is deliberately restrictive. It does not deny quantum speedup. It denies that branchwise classical computation explains it, and it denies that circuit universality should be conflated with Turing universality.

The claim is not meant to settle interpretive questions about Everettian or many-worlds accounts of quantum computation. Advocates of such accounts have argued that quantum algorithms can be understood in terms of computations distributed across branches or worlds [7–9]. The narrower point made here is operational and model-theoretic: the components of a superposition are not independently readable classical records, and the computational advantage is realized only when the global unitary evolution and final measurement context convert phase relations into an accessible statistic.

II. SUPERPOSITION IS NOT BRANCHWISE READOUT

Let an n -qubit state be written in the computational basis as

$$|\psi\rangle = \sum_{x \in \{0,1\}^n} \alpha_x |x\rangle.$$

The dimension of the Hilbert space is 2^n , and this exponential dimension is often used to motivate quantum parallelism. But the amplitudes α_x are not classical memory

* karl.svozil@tuwien.ac.at

cells. A measurement yields one outcome per run, with probabilities determined by the Born rule. Recovering the amplitudes requires many similarly prepared systems and many measurement settings, that is, tomography.

Thus a superposition can be acted on linearly, but it cannot be read branch by branch. Holevo's theorem gives one precise expression of this limitation. For an ensemble of states on a Hilbert space $\mathcal{H}$, the accessible classical information is bounded above by the Holevo quantity χ , and $\chi \leq \log_2 \dim \mathcal{H}$. Hence, for an n -qubit transmission, the accessible classical information per use is at most n bits [6, 10].

The state vector may contain exponentially many amplitudes, but the accessible classical readout is sharply constrained. This is the first reason why the naive parallelism metaphor is misleading. The superposed components can interfere, but they are not simultaneously available classical facts.

III. THE USEFUL MECHANISM: SPREAD, COMPUTE, DETECT

Many oracle-based and Fourier-sampling quantum algorithms have the schematic form

$$\begin{aligned} |0\rangle^{\otimes n} &\xrightarrow{\text{spread}} \frac{1}{\sqrt{2^n}} \sum_x |x\rangle \\ &\xrightarrow{\text{compute}} \frac{1}{\sqrt{2^n}} \sum_x (-1)^{f(x)} |x\rangle \\ &\xrightarrow{\text{detect}} \text{measurable interference pattern.} \end{aligned}$$

The middle step may be produced by phase kickback. If

$$U_f |x\rangle |y\rangle = |x\rangle |y \oplus f(x)\rangle$$

and the target is prepared in

$$|-\rangle = \frac{|0\rangle - |1\rangle}{\sqrt{2}},$$

then

$$U_f |x\rangle |-\rangle = (-1)^{f(x)} |x\rangle |-\rangle.$$

The values of $f(x)$ are not read out individually. They are converted into phases. The final measurement is useful only if those phases interfere so that the desired predicate becomes detectable.

This is the essential mechanism: spread, compute, detect. The algorithm spreads amplitude, computes coherently, and then detects a global feature through interference in a chosen measurement context.

Deutsch-Jozsa is the simplest example. After phase kickback, the query register is in the state

$$\frac{1}{\sqrt{2^n}} \sum_x (-1)^{f(x)} |x\rangle.$$

Applying the final Hadamard transform $H^{\otimes n}$ gives the amplitude of $|0^n\rangle$ as

$$\langle 0^n | H^{\otimes n} \left(\frac{1}{\sqrt{2^n}} \sum_x (-1)^{f(x)} |x\rangle \right) = \frac{1}{2^n} \sum_x (-1)^{f(x)}.$$

This amplitude is ± 1 for constant functions and 0 for balanced functions. The promise makes the global predicate visible as complete constructive or destructive interference.

Shor's algorithm is a more sophisticated example. Modular exponentiation is computed reversibly, but the algorithm does not read all values of

$$f(x) = a^x \bmod N.$$

Instead, periodicity is sampled through a quantum Fourier transform. The success comes from algebraic structure: translation symmetry becomes spectral concentration.

The same observation explains why generic predicates are not made easy by superposition. For an arbitrary Boolean string

$$b = (b_1, \dots, b_N),$$

the parity

$$b_1 \oplus b_2 \oplus \dots \oplus b_N$$

depends on every input bit. Exact quantum computation of parity requires exactly $\lceil N/2 \rceil$ queries, giving at most a factor-of-two improvement over the deterministic classical query complexity [11]. More generally, polynomial-method bounds show that for total Boolean functions the bounded-error quantum and classical query complexities are polynomially related, ruling out arbitrary exponential speedup for such black-box total functions [12]. Superposition alone is therefore not the resource. Detectable interference structure is.

IV. GARBAGE IS NOT ERASED BY UNITARITY

A second misconception concerns garbage. Classical irreversible computation often hides the physical cost of deleting information. Quantum computation cannot do so, because closed-system evolution is unitary:

$$|\psi\rangle \mapsto U |\psi\rangle, \quad U^\dagger U = I.$$

A non-injective classical function cannot be implemented by a unitary acting only on the input register. It must be embedded reversibly, for example as

$$U_f : |x\rangle |y\rangle \mapsto |x\rangle |y \oplus f(x)\rangle.$$

The input, or an equivalent record, remains present.

If a reversible computation produces

$$|x\rangle |0\rangle |0\rangle \mapsto |x\rangle |f(x)\rangle |g(x)\rangle,$$

the garbage register $|g(x)\rangle$ cannot simply be deleted. It must be uncomputed, exported, measured, or thermodynamically reset.

No isolated unitary erasure. Let $|\phi_1\rangle$ and $|\phi_2\rangle$ be two normalized pure states on a Hilbert space $\mathcal{H}$ which are not identical up to a global phase. Then there is no unitary U on $\mathcal{H}$ and no fixed normalized state $|0\rangle \in \mathcal{H}$ such that

$$U|\phi_1\rangle = U|\phi_2\rangle = |0\rangle.$$

Indeed, unitaries preserve inner products. If both equalities held, then

$$\langle\phi_1|\phi_2\rangle = \langle\phi_1|U^\dagger U|\phi_2\rangle = \langle 0|0\rangle = 1,$$

so $|\phi_1\rangle$ and $|\phi_2\rangle$ would coincide as rays, contradicting the assumption.

Consequently, if $|g(x)\rangle_B$ is a garbage register produced by reversible evaluation of a non-constant function g on input x , there is no unitary acting on B alone that resets $|g(x)\rangle_B$ to $|0\rangle_B$ for all x . The garbage can be removed only by (i) uncomputation, using the input register A , which is reversible and redistributes rather than destroys the information; (ii) coupling B to an environment, dissipating the information thermodynamically; or (iii) measurement of B , creating a classical record whose subsequent reset is subject to Landauer's principle.

Bennett's reversible-computation analysis shows how intermediate garbage can be uncomputed [13–15]; it does not make physical erasure free. Landauer's principle locates the cost in reset, coarse graining, or environmental dissipation [16]. Whether the Landauer–Bennett account fully exorcises Maxwell's demon remains a point of philosophical dispute [17–19], but the operational lesson for quantum computation is unambiguous: unitary dynamics does not delete information.

V. COPYING AND DELETION ARE CONTEXTUAL

Classical information can be copied because it is encoded in distinguishable states. Quantum information cannot generally be copied. If a unitary cloner existed with

$$U|\psi\rangle|0\rangle = |\psi\rangle|\psi\rangle \quad \text{and} \quad U|\phi\rangle|0\rangle = |\phi\rangle|\phi\rangle,$$

then unitarity preserves inner products: on the one side,

$$\langle\psi|\langle 0|U^\dagger U|\phi\rangle|0\rangle = \langle\psi|\phi\rangle\langle 0|0\rangle = \langle\psi|\phi\rangle,$$

while on the other side the cloned outputs give

$$\langle\psi\psi|\phi\phi\rangle = \langle\psi|\phi\rangle^2.$$

Equating these yields

$$\langle\psi|\phi\rangle(\langle\psi|\phi\rangle - 1) = 0,$$

so $|\psi\rangle$ and $|\phi\rangle$ are either orthogonal or coincide as rays. This is the no-cloning theorem [20, 21]. For mixed states, the no-broadcasting theorem says that a set of density operators can be broadcast if and only if they pairwise commute [22]. The no-deletion theorem gives the dual obstruction: an unknown quantum state cannot be deleted from one copy while leaving another intact [23].

These results identify the boundary of classicality. Within a fixed orthonormal basis one may copy basis labels,

$$|x\rangle|0\rangle \mapsto |x\rangle|x\rangle,$$

but this copies information only relative to that basis; it does not copy an arbitrary unknown state. In algebraic terms, classical copyability is confined to a commuting context.

Here a *context* means a maximal set of mutually commuting observables, equivalently a maximal Boolean subalgebra of the projection lattice. In the nondegenerate finite-dimensional case it corresponds to an orthonormal basis and the projectors generated by it.

Quantum computation therefore manipulates information that is not globally copyable or disposable. Only after a context is selected, typically by measurement or decoherence, do classical records emerge.

VI. CONTEXTUALITY AND THE ABSENCE OF A GLOBAL CLASSICAL HISTORY

Quantum algorithms move between incompatible descriptions: computational basis, phase basis, Fourier basis, stabilizer basis, and measurement basis. These basis changes already express complementarity. A state sharp in one basis is generally delocalized in another.

In Hilbert spaces of dimension at least three, the Kochen–Specker theorem strengthens this point: one cannot assign pre-existing values to all observables independently of the commuting context in which they are measured [24–26]. Operational versions extend contextuality to preparations, transformations, and unsharp measurements [27]. For single qubits the standard Kochen–Specker obstruction for projective measurements is absent; simple hidden-variable representations of projective qubit measurements can be constructed. This qualification is important: generalized measurements can exhibit Kochen–Specker-type contextuality even for a single qubit [28]. Nevertheless, complementarity already prevents joint sharp values for noncommuting projective observables and so blocks simultaneous classicalization of all measurement bases.

This matters for quantum computation. The intermediate stages of an algorithm cannot generally be interpreted as a single classical trajectory through pre-existing

values. The algorithm is a sequence of context changes implemented by unitary transformations and terminated by a final measurement context. What becomes classical is not the whole history, but the final record.

Resource-theoretic results support this perspective, but with important qualifications. Stabilizer operations admit efficient classical simulation by the Gottesman–Knill theorem [29, 30]; hence universal quantum advantage in the circuit model requires resources outside the efficiently simulable stabilizer fragment. Depending on the formalism and dimension, these resources may be described operationally as magic states, phase-space theoretically as Wigner-function negativity, or foundationally as contextuality [31–33]. These characterizations are related but not identical. In particular, the clean necessity results connecting contextuality with magic-state quantum computation are most direct in odd-prime-dimensional qudit settings, while the qubit case involves additional subtleties. The stable message is therefore not that contextuality alone is the complete source of speedup, but that nonclassical computational advantage requires resources beyond mere coherent superposition.

VII. QUANTUM UNIVERSALITY VERSUS TURING UNIVERSALITY

The word “universal” is used in two substantially different senses in computation theory and quantum information theory. Confusing them leads to conceptual errors.

In the quantum-circuit literature, a finite gate set is called universal if circuits over that gate set can approximate arbitrary unitary transformations on n qubits, typically to arbitrary accuracy in operator norm. In this sense, universality is a statement about the dense generation of groups such as $U(2^n)$ or $SU(2^n)$ by finite collections of gates. It is a statement about controllability of finite-dimensional Hilbert spaces.

Turing universality is different. A Turing-universal device is not merely a generator of transformations on a fixed finite-dimensional space. It acts on symbolic inputs of arbitrary length, admits arbitrarily long computations, supports parsing and composition of programs, and allows self-reference through coding of machines by strings [34, 35]. This is the setting in which partial recursive functions, undecidability, universal simulation, and the halting problem arise.

Thus a finite quantum circuit, by itself, does not define the same kind of object as a Turing machine. A circuit on n qubits implements a unitary transformation

$$U_n : \mathbb{C}^{2^n} \rightarrow \mathbb{C}^{2^n}.$$

It is a finite-dimensional, norm-preserving map. It has fixed width and depth, no unbounded tape, no open-ended input parsing, and no intrinsic recursion. It also has no autonomous mechanism for extending its own computation to arbitrary symbolic length.

A family of circuits $\{C_n\}_{n \geq 1}$ may compute a language or function in the complexity-theoretic sense only when there is a uniform rule generating the circuit C_n from n . That uniformity condition is itself algorithmic and is normally supplied classically, or by a specified higher-level quantum model such as a quantum Turing machine [7, 36, 37]. In this sense, quantum circuit families and quantum Turing machines are not identified merely by the existence of a universal gate set; their equivalence requires explicit uniformity and simulation assumptions.

This distinction is important for the phrase “quantum computable function.” In the finite-circuit model, a circuit computes only a finite input-output relation, usually probabilistically and after measurement. A uniform circuit family computes a sequence of such relations. A quantum Turing machine or another uniform quantum model can be compared to classical Turing computation, but then the usual Church–Turing issue concerns computability, while the quantum advantage concerns complexity. Standard quantum computation is not believed to enlarge the class of computable classical functions beyond the partial recursive ones; rather, it may reduce the resources needed for some computable tasks.

There is an additional physical obstruction to importing the ordinary absorbing-state picture of classical halting directly into closed unitary dynamics. A classical irreversible computation may enter an absorbing halting configuration: once halted, the machine remains halted while its output is read. This many-to-one terminal behavior is incompatible with isolated unitary evolution. The point is not that standard quantum models lack operational termination; quantum Turing machines and circuit models use halt qubits, clocks, measurement protocols, or externally prescribed stopping rules [7, 36, 38, 39]. The point is narrower: a closed unitary cannot implement a nontrivial absorbing halt state into which distinct pre-halting states collapse.

Indeed, suppose a normalized halting state $|h\rangle$ is fixed,

$$U|h\rangle = |h\rangle,$$

and suppose some distinct normalized pre-halting state $|\psi\rangle \neq e^{i\theta}|h\rangle$ also evolves into the same halting state,

$$U|\psi\rangle = |h\rangle.$$

Then

$$U|\psi\rangle = U|h\rangle,$$

and since U is unitary, hence invertible, hence injective, this forces $|\psi\rangle = |h\rangle$, contradicting the assumption. Thus a unitary may coherently mark, clock, or correlate halting information, but it cannot many-to-one collapse distinct computational histories into a single absorbing terminal configuration without exporting information to additional degrees of freedom.

This observation should not be confused with a no-go theorem against quantum computation or quantum Turing machines. Myers emphasized difficulties associated

with fully quantum halting schemes, while Ozawa and subsequent analyses refined monitoring and halt-flag protocols so that operational halting can be treated without spoiling the computation [38, 39]. The present claim is only that the classical absorbing halt state is not itself a closed-unitary primitive.

This reinforces the main thesis. Quantum circuit universality is a powerful notion of finite-dimensional controllability and approximation. Turing universality is a notion of symbolic, unbounded, self-referential computation under an explicit uniformity condition. They intersect in uniform models of quantum computation, but they are not identical. Treating “universal quantum computation” as if it automatically meant Turing universality obscures the distinction between unitary dynamics, measurement, and algorithmic control.

VIII. MINIATURIZATION AND HIGH-DIMENSIONAL REPRESENTATION

There are two broader motivations for this reframing.

First, miniaturization makes quantum information processing unavoidable. Classical bits are physical metastable systems. Their stability depends on energy barriers, redundancy, dissipation, and environmental decoherence. As devices shrink, discreteness, tunneling, thermal fluctuations, measurement back-action, and coherence become harder to ignore. Classical computation is not the fundamental regime; it is an engineered regime stabilized by einselection through environmental decoherence [40]. Quantum computation may therefore be less an exotic alternative than the natural information theory of sufficiently small controlled systems.

Second, there is a limited but useful analogy with high-dimensional representation learning. Large vector spaces allow many directions to be nearly orthogonal, and concentration of measure makes such quasi-orthogonality typical in high dimension [41, 42]. Quantum Hilbert spaces likewise exploit high-dimensional geometry. But in neither case does large representational dimension imply unrestricted readout. Information must be extracted through projections, measurements, sampling, or compression. Thus high dimension is a resource for structured representation and interference, not a license for arbitrary classical access.

The shared moral is:

large dimension $\nRightarrow$ unlimited readout.

In quantum computation, this limitation is enforced by measurement, Holevo bounds, no-cloning, and contextuality. In vector semantics, it appears as finite context, finite sampling, and the need to compress distributed representations into discrete outputs.

IX. CONCLUSION

Quantum computation is not classical computation performed in exponentially many parallel branches whose results can later be inspected. Its power lies in reversible amplitude manipulation, phase-sensitive computation, and contextual measurement. Speedup appears when the desired predicate can be encoded into a unitary process and made visible by interference in an appropriate final context.

The negative constraints are just as important as the positive mechanism. Garbage cannot be erased by unitary dynamics; it must be uncomputed or exported, and physical reset is thermodynamic. Unknown quantum states cannot be freely copied, broadcast, or deleted. Contextuality prevents the assignment of a single global classical history to the computation.

Nor should circuit or unitary universality be confused with Turing universality. A finite universal gate set provides dense control over unitary transformations; it is not, by itself, a symbolic machine with arbitrary parsing, recursion, self-reference, and a classical halting semantics. In closed unitary dynamics there is no nontrivial absorbing halt state of the classical many-to-one kind. Termination and output extraction require clocks, flags, measurements, open-system records, or external control.

Thus quantum computation is neither magic parallelism nor ordinary classical computation in disguise. It is context-dependent, thermodynamically embedded, reversible information processing whose useful answers must be engineered as interference phenomena.

ACKNOWLEDGMENTS

This research was funded in whole or in part by the *Austrian Science Fund (FWF)* [Grant DOI:10.55776/PIN5424624]. The author acknowledges TU Wien Bibliothek for financial support through its Open Access Funding Programme.

This text was partially created and revised with assistance from large language models. All content, ideas, and prompts were provided by the author.

-
- [1] D. Deutsch and R. Jozsa, Rapid solution of problems by quantum computation, *Proceedings of the Royal Society: Mathematical and Physical Sciences* (1990-1995) **439**, 553 (1992).
 - [2] D. R. Simon, On the power of quantum computation, *SIAM Journal on Computing* **26**, 1474 (1997).

- [3] P. W. Shor, Algorithms for quantum computation: discrete logarithms and factoring, in *Proceedings of the 35th Annual Symposium of on Foundations of Computer Science, Santa Fe, NM, Nov. 20-22, 1994* (IEEE Computer Society Press, 1994) [arXiv:quant-ph/9508027](https://arxiv.org/abs/quant-ph/9508027).
- [4] L. K. Grover, A fast quantum mechanical algorithm for

- database search, in *Proceedings of the Twenty-Eighth Annual ACM Symposium on the Theory of Computing* (1996) pp. 212–219, [arXiv:quant-ph/9605043](#).
- [5] R. Cleve, A. Ekert, C. Macchiavello, and M. Mosca, Rapid solution of problems by quantum computation, *Proceedings of the Royal Society A: Mathematical, Physical and Engineering Sciences* **454**, 339 (1998).
 - [6] M. A. Nielsen and I. L. Chuang, *Quantum Computation and Quantum Information* (Cambridge University Press, Cambridge, 2000).
 - [7] D. Deutsch, Quantum theory, the Church-Turing principle and the universal quantum computer, *Proceedings of the Royal Society of London. Series A, Mathematical and Physical Sciences* (1934-1990) **400**, 97 (1985).
 - [8] C. Hewitt-Horsman, An introduction to many worlds in quantum computation, *Foundations of Physics* **39**, 869 (2009).
 - [9] D. Wallace, *The Emergent Multiverse: Quantum Theory according to the Everett Interpretation* (Oxford University Press, 2012).
 - [10] A. S. Holevo, Bounds for the quantity of information transmitted by a quantum communication channel, *Problemy Peredachi Informatsii* **9**, 3 (1973), english translation in *Problems of Information Transmission* **9**, 177–183 (1973).
 - [11] E. Farhi, J. Goldstone, S. Gutmann, and M. Sipser, Limit on the speed of quantum computation in determining parity, *Physical Review Letters* **81**, 5442 (1998), [arXiv:quant-ph/9802045](#).
 - [12] R. Beals, H. Buhrman, R. Cleve, M. Mosca, and R. de Wolf, Quantum lower bounds by polynomials, *Journal of the ACM* **48**, 778 (2001), [arXiv:quant-ph/9802049](#).
 - [13] C. H. Bennett, Logical reversibility of computation, *IBM Journal of Research and Development* **17**, 525 (1973).
 - [14] C. H. Bennett, The thermodynamics of computation—a review, *International Journal of Theoretical Physics* **21**, 905 (1982).
 - [15] C. H. Bennett, Notes on Landauer’s principle, reversible computation, and Maxwell’s demon, *Studies in History and Philosophy of Science Part B: Studies in History and Philosophy of Modern Physics* **34**, 501 (2003).
 - [16] R. Landauer, Irreversibility and heat generation in the computing process, *IBM Journal of Research and Development* **5**, 183 (1961).
 - [17] J. Earman and J. D. Norton, EXORCIST XIV: The wrath of Maxwell’s demon. part I. from Maxwell to Szilard, *Studies In History and Philosophy of Science Part B: Studies In History and Philosophy of Modern Physics* **29**, 435 (1998).
 - [18] J. Earman and J. D. Norton, EXORCIST XIV: The wrath of Maxwell’s demon. part II. from Szilard to Landauer and beyond, *Studies In History and Philosophy of Science Part B: Studies In History and Philosophy of Modern Physics* **30**, 1 (1999).
 - [19] J. D. Norton, Eaters of the lotus: Landauer’s principle and the return of Maxwell’s demon, *Studies in History and Philosophy of Science Part B: Studies in History and Philosophy of Modern Physics* **36**, 375 (2005).
 - [20] W. K. Wootters and W. H. Zurek, A single quantum cannot be cloned, *Nature* **299**, 802 (1982).
 - [21] D. Dieks, Communication by EPR devices, *Physics Letters* **92A**, 271 (1982).
 - [22] H. Barnum, C. M. Caves, C. A. Fuchs, R. Jozsa, and B. Schumacher, Noncommuting mixed states cannot be broadcast, *Physical Review Letters* **76**, 2818 (1996).
 - [23] A. K. Pati and S. L. Braunstein, Impossibility of deleting an unknown quantum state, *Nature* **404**, 164 (2000).
 - [24] S. Kochen and E. P. Specker, The problem of hidden variables in quantum mechanics, *Journal of Mathematics and Mechanics (now Indiana University Mathematics Journal)* **17**, 59 (1967).
 - [25] J. S. Bell, On the problem of hidden variables in quantum mechanics, *Reviews of Modern Physics* **38**, 447 (1966).
 - [26] D. N. Mermin, Hidden variables and the two theorems of John Bell, *Reviews of Modern Physics* **65**, 803 (1993).
 - [27] R. W. Spekkens, Contextuality for preparations, transformations, and unsharp measurements, *Physical Review A* **71**, 052108 (2005), [arXiv:quant-ph/0406166](#).
 - [28] A. Cabello, Kochen–specker theorem for a single qubit using positive operator-valued measures, *Physical Review Letters* **90**, 190401 (2003).
 - [29] D. Gottesman, *The Heisenberg representation of quantum computers* (1998), [arXiv:quant-ph/9807006](#).
 - [30] S. Aaronson and D. Gottesman, Improved simulation of stabilizer circuits, *Physical Review A* **70**, 052328 (2004).
 - [31] M. Howard, J. Wallman, V. Veitch, and J. Emerson, Contextuality supplies the ‘magic’ for quantum computation, *Nature* **510**, 351 (2014).
 - [32] V. Veitch, S. A. H. Mousavian, D. Gottesman, and J. Emerson, The resource theory of stabilizer quantum computation, *New Journal of Physics* **16**, 013009 (2014).
 - [33] R. Raussendorf, Contextuality in measurement-based quantum computation, *Physical Review A* **88**, 022322 (2013).
 - [34] A. M. Turing, On computable numbers, with an application to the Entscheidungsproblem, *Proceedings of the London Mathematical Society* **2**, **42**, 230 (1937).
 - [35] A. M. Turing, On computable numbers, with an application to the Entscheidungsproblem. A correction, *Proceedings of the London Mathematical Society* **2**, **43**, 544 (1938).
 - [36] E. Bernstein and U. Vazirani, Quantum complexity theory, *SIAM Journal on Computing* **26**, 1411 (1997), preliminary version in *Proceedings of the 25th Annual ACM Symposium on Theory of Computing* (STOC 1993), pp. 11–20.
 - [37] H. Nishimura and M. Ozawa, Computational complexity of uniform quantum circuit families and quantum Turing machines, *Theoretical Computer Science* **276**, 147 (2002).
 - [38] J. M. Myers, Can a universal quantum computer be fully quantum?, *Physical Review Letters* **78**, 1823 (1997).
 - [39] M. Ozawa, Quantum nondemolition monitoring of universal quantum computers, *Physical Review Letters* **80**, 631 (1998).
 - [40] W. H. Zurek, Decoherence, einselection, and the quantum origins of the classical, *Reviews of Modern Physics* **75**, 715 (2003).
 - [41] M. Ledoux, *The Concentration of Measure Phenomenon* (American Mathematical Society, 2001).
 - [42] R. Vershynin, *High-Dimensional Probability: An Introduction with Applications in Data Science* (Cambridge University Press, 2018).