

Computable Countermarkets and the Limits of Universal Trading

Karl Svozil ¹

¹*Institute for Theoretical Physics, TU Wien, Wiedner Hauptstraße 8-10/136, A-1040 Vienna, Austria**
(Dated: September 19, 2026)

We explain why no trading algorithm can guarantee profit in every market. For each deterministic program that always returns a finite-precision position, we construct a fixed, algorithmically generated price path on which every active position loses and inactivity earns nothing. This holds with positive, continually changing prices, costless trading, and unlimited computation time. Separate arguments limit learning market rules, certifying future events, and establishing randomness from finite data; financial assumptions excluding risk-free gains without initial investment impose another boundary. Useful strategies may exploit market structure, information, or compensation for risk, while benchmark performance need not imply profit. Reversing and rearranging price histories within the assumed market class provide practical stress tests, distinguishing conditional success from universal guarantees.

I. INTRODUCTION

Throughout this paper the elementary configuration has two sides: a *trader* and a *market*. At each date the trader converts observed history into a position, meaning its chosen asset holdings; the market supplies the next price. Depending on the question, the market side is a realized path (a sequence of prices), a path-generating program, a probability law assigning probabilities to possible outcomes, or a price process constrained by no-arbitrage. A price process models prices as random variables over time; no-arbitrage excludes a risk-free opportunity to gain from zero initial capital. Holding the trader fixed while changing the admitted market side reveals which universal claim is at issue. A pathwise claim concerns individual price paths, rather than averages over a probability law.

A strategy is *universal* here only if, after discounting by a specified numeraire, it produces strictly positive terminal gain from zero initial capital on every admitted trajectory at a specified finite horizon. The numeraire is a positive-valued reference asset, here a bank account; discounting expresses prices and gains in units of that asset. The horizon is the final evaluation date, and a trajectory is a price path. This claim differs from nonnegative gain everywhere with positive gain somewhere, eventual profit at a path-dependent date, expected profit (a probability-weighted average), and small regret to a benchmark. A benchmark is a specified reference investment or strategy; regret measures a strategy’s performance shortfall relative to it. Zero cost means zero initial investment, self-financing means no later external cash injections, and admissibility imposes restrictions on allowable trading, such as a lower bound on wealth. These conditions distinguish trading gain from positive wealth obtained by investing positive capital. The order of the quantifiers, which specify “there exists” and “for every,” is

$$\exists \sigma \forall \mathcal{E} : \quad \sigma \text{ wins against } \mathcal{E}. \quad (1)$$

Here σ denotes a trading strategy and $\mathcal{E}$ an admitted market; $\exists$ means “there exists” and $\forall$ means “for every.” “Wins” means earning strictly positive discounted gain from zero initial capital at the specified finite horizon. The order requires one fixed strategy to succeed against every admitted market. By contrast, diagonalization—constructing a counterexample tailored to the proposed strategy—establishes

$$\forall \sigma \exists \mathcal{E}_\sigma : \quad \sigma \text{ does not win against } \mathcal{E}_\sigma. \quad (2)$$

The subscript in $\mathcal{E}_\sigma$ indicates that the countermarket may depend on the chosen strategy; there need not be one market defeating all strategies.

A computable rule is one executable by an algorithm; “effective” and “recursive” have this same computational meaning here. A total rule returns an answer on every allowed input, and a deterministic rule makes no random choices. For a fixed total deterministic computable trader, its code can be embedded in a total program that recursively chooses the next bounded proportional price move opposite to each nonzero position and a nonzero move during inactivity. The resulting countermarket is one fixed computable path; any finite prefix can be generated offline. A trader that never takes a position earns zero, which already refutes a guarantee of strict profit. This is the trading counterpart of the next-bit construction $x_0 = 1 - F(\emptyset)$ and $x_t = 1 - F(x_0, \dots, x_{t-1})$ for $t \geq 1$: it defeats the fixed proposed rule, not every rule on one exceptional sequence. Here a bit is either zero or one, F predicts the next bit from the preceding bits, $\emptyset$ denotes the empty history, and x_t is the bit chosen opposite to that prediction at date t . A prefix is an initial finite segment of a sequence; generating it offline means computing it without interacting with a live market.

Three complementary limits use different market information. Gold’s learner sees an accumulating history and must eventually stabilize on a correct program; no effective learner identifies every binary total recursive time function (a program assigning a bit to every date), although restricted classes such as the primitive-recursive time functions are identifiable in the limit [1, 2].

* karl.svozil@tuwien.ac.at; <http://tph.tuwien.ac.at/~svozil>

Primitive-recursive functions are built from elementary functions using composition and recursion whose number of stages is set by the input. Identification in the limit means eventually keeping one correct rule, without necessarily knowing when that point has been reached. A source-code procedure that settled every future event of Turing-universal generators, which can simulate any program, would instead decide the Halting Problem: whether an arbitrary program eventually stops. Busy-Beaver growth, the growth of the longest halting runtime among programs up to a given length, excludes a computable worst-case waiting horizon. Finally, relative to a specified computable probability law, a Martin-Löf-random record escapes every effective null test: an algorithmically specified test whose full failure set has probability zero. Such records have measure one, meaning probability one, also called “almost surely.” Finite samples cannot certify this infinite-sequence property. No-arbitrage supplies a separate boundary through a pricing measure, a probability law used to value discounted pay-offs.

Repeatable practical success must therefore identify structure absent from the unrestricted market class: persistence (continuation of past behavior), mean reversion (movement back toward a reference level), a risk premium (expected compensation for bearing risk), a benchmark, lawful information, or another distributional asymmetry (an imbalance in possible outcomes or their probabilities). An edge is an expected performance advantage under specified market conditions. Large, crowded, or public strategies may also change the distribution through price impact, meaning price changes caused by their own orders, and adaptation by other participants. This falsifiable restriction, one that evidence could contradict, is the operational meaning of “insight” used below.

The explicit diagonal countermarket is the constructive core of the paper. Table I organizes the complementary limits into a taxonomy that keeps distinct several statements commonly summarized as “no universal strategy.” The construction and its proofs are expository reformulations rather than a claim of a newly discovered impossibility theorem. The analysis connects this quantifier map to the literature on unpredictable prices [3–5], Gold learning, computability and randomness, no-arbitrage and no-free-lunch (NFL) results about performance averaged across possible problems, time-reversal and permutation stress testing (reversing or otherwise reordering a record), an options-strategy case study, Cover universality, and conditional strategies that work “for all practical purposes” (FAPP).

II. THE TRADER AND THE MARKET: COMPUTABILITY AND RANDOMNESS

We now formalize the trader–market pair. The trader is a program M_σ mapping each observed finite price history to a position; the market side supplies the price

history. Four specifications lead to four different boundaries: a fixed computable countermarket tailored to a trader defeats its claim of pathwise universality; Gold’s history protocol precludes identifying every computable rule in the limit; a computationally universal class of market generators makes uniform certification of unbounded future events undecidable, meaning that no algorithm can always give a correct yes-or-no answer and terminate; and a computable probability law links process randomness to individual random records through an almost-sure statement.

A. A Computable Countermarket for Every Trader

Throughout the discrete constructions, dates $t = 0, 1, 2, \dots$ count trading periods, and T denotes a finite final date. The quantity $P_t = S_t/B_t$ is the discounted price of one asset, where S_t is its price in a fixed currency and $B_t > 0$ is the value of one unit of the chosen risk-free bank account. For example, if both initially equal 100, and later $S_1 = 121$ and $B_1 = 110$, the discounted price rises from 1 to 1.1. This measures performance relative to the bank account, not inflation. The position w_t is a number of shares chosen before P_{t+1} is observed. A positive position holds shares, a negative position sells borrowed shares short, and zero means no asset exposure. A short position profits from a price fall and loses from a rise. Trading is self-financing: changes in holdings are funded through the existing share and bank balances, with no external cash injections. Let V_t be total portfolio wealth at date t , measured in bank-account units. More explicitly, holding w_t shares and the residual $V_t - w_t P_t$ in the discounted bank account gives

$$V_{t+1} = V_t + w_t(P_{t+1} - P_t).$$

We omit transaction costs, borrowing limits, leverage caps (limits on exposure relative to capital), and solvency constraints (requirements to remain able to meet obligations). The counterexample therefore also applies to any subclass of these strategies satisfying additional restrictions.

Let $M_\sigma: \mathbb{Q}_{>0}^{<\omega} \rightarrow \mathbb{Q}$ be a total computable function on finite rational histories, encoded in the usual effective way, with $w_t = M_\sigma(P_0, \dots, P_t)$. Here $\mathbb{Q}$ is the set of rational numbers, $\mathbb{Q}_{>0}$ its positive members, and the superscript $< \omega$ denotes all finite sequences; the arrow specifies the function’s input and output sets. Totality requires an action in finite computation time on every input; it supplies no prescribed running-time bound or guarantee of meeting a trading deadline. Rational outputs model finite-precision orders and, crucially, make comparison with zero decidable. Let $\mathcal{P}_{\text{comp}}^+$ denote the positive rational-valued computable price paths.

TABLE I. Distinct claims organized by the trader–market configuration. They are complementary, not interchangeable. Here P_σ is the price path constructed for strategy σ ; a black box reveals outputs but not its internal program, and a promise restricts the inputs on which an algorithm must work.

Mechanism	Market side	Formal scope	Conclusion
Diagonal countermarket	Fixed total program constructed from the fixed trader’s code	Pathwise worst case: $\forall \sigma \exists P_\sigma$	No total deterministic computable trader wins on every computable price path
Gold identification in the limit	Accumulating computable history, or interaction with a black box	Success is eventual stabilization on a correct program, without a known stopping date	No effective learner identifies every computable binary time function; restricted classes can be identified
Halting reduction and Busy Beaver	Total program or generator, under an explicit promise	Individual encoded computation over an unbounded horizon	No general computable profit or market-event decider, or computable description-size waiting horizon
Algorithmic randomness	Passive binary record under the computable fair-coin reference law	Each Martin-Löf-random path; such paths have measure one	Independent fair trials yield Martin-Löf-random records almost surely; finite tests cannot certify randomness
No-arbitrage	Stochastic price process admitting a coherent pricing measure	Payoff conditions holding with probability one under a pricing measure	No admissible zero-cost classical arbitrage
Cover universality	Passive bounded sequence and a stated benchmark class	Pathwise and benchmark-relative as the horizon grows	Average regret per period tends to zero, without guaranteeing positive profit

Theorem 1 (Computable Diagonal Countermarket). *Fix a rational initial price $P_0 > 0$ and a rational proportional price-change magnitude $0 < \epsilon < 1$. For every total deterministic computable strategy M_σ , there exists a fixed path $P^{M_\sigma} \in \mathcal{P}_{\text{comp}}^+$, starting at P_0 , with $P_{t+1}/P_t \in \{1 - \epsilon, 1 + \epsilon\}$ at every date, on which its self-financing gain*

$$G_T = \sum_{t=0}^{T-1} w_t(P_{t+1} - P_t)$$

is non-positive at every finite horizon T . It is strictly negative after the strategy has taken a nonzero position at least once.

The superscript in P^{M_σ} labels dependence on the trader’s program; it is not a power. The gain G_T sums the position times the price change over the first T periods and equals final discounted wealth when $V_0 = 0$. The proportional return is the price change divided by the preceding price; its magnitude here is always ϵ .

Proof. Choose rational $P_0 > 0$ and rational $0 < \epsilon < 1$. Recursively compute $w_t = M_\sigma(P_0, \dots, P_t)$ and set

$$P_{t+1} = \begin{cases} P_t(1 - \epsilon), & w_t > 0, \\ P_t(1 + \epsilon), & w_t < 0, \\ P_t(1 + \epsilon(-1)^t), & w_t = 0. \end{cases} \quad (3)$$

Because M_σ is total and rational-valued, this recursion is a total program for one fixed infinite price path. Its

prices remain positive and computable, and every return has magnitude ϵ . In each period, $w_t(P_{t+1} - P_t) \leq 0$, with strict inequality when $w_t \neq 0$. Summing proves the claim at every finite horizon. $\square$

In the last case of Eq. (3), the alternating market move is a convenient but arbitrary choice. Since $w_t = 0$, any market move yields zero gain in that period. Within the theorem’s fixed-magnitude class, any computable choice between returns $+\epsilon$ and $-\epsilon$ is admissible, allowing a wide variety of market-move sequences in place of the alternating convention.

For a concrete example, fix the rule that holds one share for the first period, shorts one share for the second, and then holds no shares. With $P_0 = 100$ and $\epsilon = 1/10$, Eq. (3) produces

$$(w_0, w_1, w_2) = (1, -1, 0), \\ (P_0, P_1, P_2, P_3) = (100, 90, 99, 108.9).$$

The successive discounted gains are $-10, -9, 0$, so $G_3 = -19$, all in bank-account units. The prices are constructed after fixing the rule; they are illustrative, not market data.

Remark 1 (Logical and Economic Scope of the Countermarket). Theorem 1 uses the trader’s *code*, not a physical market’s observation of a deployed order. Once M_σ is fixed, the countermarket program is fixed and can generate any desired finite prefix offline. It is therefore a lawful, deterministic, computable counterexample, although not

necessarily an efficient one computationally. If the market side is regarded as an opponent whose objective is to prevent this trader's gain, the rule is rational relative to that objective. It is not thereby shown to be an economic equilibrium, with mutually consistent participant choices, or an arbitrage-free price process, and another trader may profit on it. Such a trader has, in turn, its own computable countermarket.

This is a direct strategy-relative diagonal construction. The proposed universal trader is used to compute an object outside its success set: long is paired with down, short with up, and neutrality with an alternating nonzero return. Just as the complement sequence defeats a purported predictor of every computable sequence, P^{M_σ} defeats the purported trader on every horizon. Enumeration of all programs is unnecessary because the universal claim itself supplies the program to be diagonalized against. The quantifiers are $\forall M_\sigma \exists P^{M_\sigma}$. A constant discounted price path already gives every self-financing rule zero trading gain and refutes strict profit on a class containing that path. The stronger content of Theorem 1 is that every nonzero position of the selected trader incurs a strict loss even on the class with fixed return magnitude ϵ , which contains no flat path. It does not supply a path on which every strategy strictly loses: permanently long and permanently short positions cannot both lose on the same one-step price move.

B. Learning a Market Rule from History: Gold's Limit

A trader normally receives neither the market generator's source code nor a proof of its law. The trader sees an expanding record $x_0, x_1, \dots$ and must infer from each finite prefix how it will continue. Gold's function-identification model formalizes exactly this information flow. At time n , an effective learner L , itself an algorithm, maps $(x_0, \dots, x_n)$ to the index, or numbered program description, of a Turing machine, an abstract model of computation. It *identifies x in the limit* if, from some finite date onward, every guess is the same index and that machine computes the whole sequence. The learner need not know when its last change of mind has occurred [1, 2].

Let REC_2 denote the total computable binary sequences. Gold's negative theorem for binary-valued time functions has the trader–market quantifier order, with L ranging over effective learners,

$$\begin{aligned} \forall L \exists x \in \text{REC}_2 : \\ L(x_0, \dots, x_n) \text{ does not stabilize} \\ \text{on a program for } x. \end{aligned} \quad (4)$$

The adverse history depends on the learner. This is closely analogous to $\forall M_\sigma \exists P^{M_\sigma}$, although Gold's failure criterion is failure to converge to a fixed correct program name, not financial loss.

The corresponding claim about literal next-move forecasting is sharper and has the short direct diagonal proof already used above. If a total computable predictor F outputs one bit from every finite binary history, the recursive definition

$$x_0 = 1 - F(\emptyset), \quad x_t = 1 - F(x_0, \dots, x_{t-1}) \quad (t \geq 1)$$

produces a computable market record on which F is wrong at every date. Thus the proposed statement—every total next-move predictor has a computable countermarket—is correct, but it is this direct prediction diagonalization, not the literal statement of Gold's identification theorem. Nonconvergence of program names alone does not count the learner's next-step errors.

The opposite-bit construction also appears explicitly in Legg's analysis of universal prediction [6]. The qualifier *computable* is essential. Solomonoff induction combines program-based predictive models, giving shorter descriptions greater weight. It defines an uncomputable universal predictive ideal; for computable probabilistic environments its conditional predictions converge, with quantitative error bounds developed by Hutter [7–10]. It therefore evades, rather than contradicts, the diagonal result: it is not a total executable trading rule, and predictive convergence is not Gold's requirement of stabilizing on one exact generating program. For probability-valued forecasts, V'yugin gives a further, distinct impossibility: a randomized construction can make any partial computable forecaster fail calibration with probability arbitrarily close to one [11]. A partial algorithm may fail to return an answer on some inputs; calibration requires forecast probabilities to agree with observed frequencies on the corresponding forecast groups. Calibration, next-bit error, and identification in the limit are three different criteria.

Gold's results also show why the restriction on the market side matters. Every uniformly effective indexed family of total computable functions—one algorithm computes any member from its index and input—is identifiable in the limit: a learner selects the least index consistent with the observed prefix, and eventually eliminates every smaller incorrect index. In particular, primitive-recursive time functions form such a family. In Gold's interactive black-box model, where the eventual hypothesis must reproduce future input–output behavior, finite-state boxes, with only finitely many internal states, are weakly identifiable in the limit, whereas the class of primitive-recursive black boxes is not [1]. These are different object classes, not contradictory claims. Gold's 1967 language model further makes explicit that learnability depends on the admitted class, the presentation of the evidence, and the allowed names [2]. For trading, a forecast can therefore be justified only relative to a specified restriction on possible markets.

C. Universal Computation and the Halting Barrier

Gold’s learner works from observations and may revise its hypothesis indefinitely. The next results ask a separate, stronger operational question: can one terminating program inspect arbitrary encoded traders or market generators and correctly decide their unbounded future behavior? The preceding diagonal construction already disproves a trader that wins on every computable path and does not invoke the Halting Problem. The source-code event-certification question does.

Theorem 2 (Undecidability of Eventual Profit). *Fix the positive computable price path $P_t = t + 1$, $t = 0, 1, \dots$. There is no algorithm which, under the promise that its input code computes a total trading program M_σ , correctly decides for every such code whether*

$$\exists T \geq 1 : \quad G_T(M_\sigma; P) = \sum_{t=0}^{T-1} w_t(P_{t+1} - P_t) > 0.$$

Proof. Assume such a profit certifier exists. Given an arbitrary Turing machine A and input x , construct a trading program $M_{A,x}$ as follows. At date t , simulate $A(x)$ for $t + 1$ steps. If it has halted within those steps, set $w_t = 1$; otherwise set $w_t = 0$. Every date requires only a finite simulation, so $M_{A,x}$ is total. Along $P_t = t + 1$, every long position earns one unit in the next period. Hence $G_T > 0$ for some T if and only if $A(x)$ eventually halts. A certifier for eventual trading profit would therefore decide the Halting Problem, which is impossible [12]. $\square$

This reduction, a conversion of one decision problem into another that preserves its answer, can also be placed directly on the market side, which is the form most closely connected to source-code event certification.

Theorem 3 (Undecidability of a Market Event). *There is no algorithm which, under the promise that its input code computes a total positive-price generator G , correctly decides for every such code whether its generated path P^G ever satisfies $P_t^G > P_0^G$.*

Proof. Given a Turing machine A and input x , define the total generator by $P_0^{A,x} = 1$ and, for $t \geq 1$,

$$P_t^{A,x} = 1 + \mathbf{1}\{A(x) \text{ halts within } t \text{ steps}\}, \quad t \geq 1. \quad (5)$$

Here $\mathbf{1}\{\cdot\}$ is the indicator: one when its condition holds and zero otherwise. Each quoted price requires only a finite t -step simulation, so the generator is total and computable. Its initial price is one, and it ever quotes a higher price if and only if $A(x)$ halts. A universal decider for the stated market event would therefore decide the Halting Problem. $\square$

Theorem 3 makes “universal computation” precise: the admissible generators can emulate an arbitrary machine and expose its outcome through a price event. No algorithm can then certify every future event of every

such market. Both reductions concern promise domains: the putative decider must halt and be correct on every promised total input. Positive instances are semidecidable: simulation confirms them when the event occurs, but may run forever when it does not. There is no terminating general procedure for certifying its perpetual absence.

This is not Gold’s theorem. The reduction receives source code and must halt with a yes-or-no answer; Gold’s learner receives accumulating data and need only stabilize eventually. Indeed, the paths in Eq. (5) are identifiable in the limit: guess the flat path, then switch permanently if the price becomes two. No finite date, however, certifies that the initial guess is final. Thus halting excludes uniform certification, whereas Gold excludes identification in the limit for the class of all binary total recursive time functions. Restricted model classes may still be learnable [13–17].

a. Busy-Beaver horizons: finite glimpses and non-computable waiting Fix a step-counted universal machine V with an effective compiler that adds only fixed overhead for a given program schema, or template. The compiler translates descriptions into code for V ; fixed overhead is a constant number of extra code bits. Define the runtime Busy Beaver function by

$$\text{BB}(n) = \max\{\text{time}_V(p) : |p| \leq n \text{ and } V(p) \text{ halts}\}, \quad (6)$$

with value zero if the set is empty. Here p is a program, $|p|$ its length in bits, and $\text{time}_V(p)$ its number of execution steps on V . The input n is a program-description budget, not a date; $\text{BB}(n)$ is a runtime, which becomes a possible trading date when the computation is embedded in a market generator. A small increase in rule length may therefore require waiting beyond every computably prescribed scale.

Each $\text{BB}(n)$ is finite, and selected low values can sometimes be proved by exhaustive simulation plus proofs that the remaining programs do not halt. This is not uniform “computability for small arguments”: any finite table can be hard-coded, and its apparent pattern licenses no extrapolation. If a total computable h upper-bounded BB , simulating every program of length at most n for $h(n)$ steps would decide halting. Eventual domination uses a further program-schema argument. For each n , a program of $O(\log n)$ bits—at most a constant times $\log n$ for sufficiently large n —can encode n , compute $h(n)$, execute more than $h(n)$ computation steps solely to delay termination, and halt. Here and below, $\log$ denotes the natural logarithm. For all sufficiently large n , its compiled V -code has length at most n , so $\text{BB}(n) > h(n)$. Thus, up to the fixed machine convention, BB eventually exceeds every total computable function [18, 19].

Calude, Dinneen, and Shu illustrate the finite/infinite divide. For one explicit universal machine U with prefix-free program codes, meaning no valid code is a proper prefix of another, computation and mathematical proof settled halting through 84-bit programs and certified 64 initial bits of its halting probability Ω_U . This is the prob-

ability that U halts when its program bits are supplied by independent fair coin tosses. They did not compute Busy Beaver values, and the work needed to recover the relevant halting programs from longer supplied prefixes of Ω_U outruns every computable bound. Their result is a formidable finite glimpse, not a scalable computation of an uncomputable object [20].

Applied to Eq. (5), a short generator can quote $P_t = 1$ for a Busy-Beaver-scale interval and then quote $P_t = 2$. Across encoded programs of length at most n , the last halt can occur at $\text{BB}(n)$, up to fixed wrapper overhead. Knowing that horizon would make continued silence a certificate of no later move. Equivalently, with TotComp denoting total computable size-dependent schedules,

$$\forall h \in \text{TotComp} \exists p : V(p) \text{ halts,} \quad (7)$$

$$\text{time}_V(p) > h(|p|).$$

If c is the wrapper's fixed description-length overhead, it is absorbed by the computable envelope $H(n) = \max_{k \leq n+c} h(k)$. Every quote at a specified finite date remains computable by bounded simulation; the generator neither computes Busy Beaver nor selects its maximizer uniformly. Noncomputability lies in certification across the family. Each realized path is deterministic and computable, hence not Martin-Löf random under fair-coin coding.

This barrier is stronger than sensitive dependence on initial conditions or visual complexity alone in a precise decision-theoretic sense. Repeatedly applying finitely many distance-shrinking maps, or a chaotic map in which nearby initial states can separate strongly, can generate intricate output from a short computable rule. Turing universality additionally creates questions about the unbounded future for which no terminating general algorithm exists. The categories can overlap when a dynamical system performs universal computation. This is a worst-case claim across generator descriptions, not evidence that typical financial markets exhibit Busy-Beaver delays [17, 21, 22].

These obstructions concern unbounded horizons. At any specified finite horizon, total programs can still be run and their realized prices and gains computed.

D. Worst-Case Value and Private Randomization

Returning to the direct countermarket, its finite-horizon value can be stated explicitly, and the role of a trader's hidden randomization can be separated from deterministic pathwise opposition. The maximin value is the best worst-case lower bound on gain: the inner infimum ranges over markets, and the outer supremum ranges over strategies.

Corollary 4 (Maximin Value). *At every finite horizon T , if the neutral strategy (always holding zero shares) is allowed, then*

$$\sup_{\text{total computable } M_\sigma} \inf_{P \in \mathcal{P}_{\text{comp}}^+} G_T(M_\sigma, P) = 0.$$

Proof. The neutral strategy guarantees zero, while Theorem 1 gives every strategy a computable countermarket with gain at most zero. $\square$

a. Two sources of randomness Randomness can enter on either side: *private trader randomization* randomly selects positions, while *market randomness* generates price movements. They can occur separately or together and need not have the same effect on profits. Private randomization need not destroy an existing edge, and random market movements can have positive expected discounted returns. The result below uses the stronger assumption that the next discounted price change has conditional mean zero, even given the trader's private seed, the random input used to choose positions. Thus the expected next change, given all available information, is zero. Integrability means finite expected absolute value; under this condition the gain has expectation zero at each fixed horizon, not convergence of realized profits to zero in the long run.

Proposition 5 (Private Randomization Does Not Create a Guarantee). *Fix a finite horizon T , a rational $P_0 > 0$, and a rational $\epsilon \in (0, 1)$. There is a single passive computable market law, independent of the trader, under which every predictable randomized strategy, choosing positions using only information already available, with integrable gains has expected gain zero. Specifically, let*

$$P_{t+1} = P_t(1 + \epsilon \xi_{t+1}), \quad t = 0, \dots, T-1,$$

where $(\xi_t)_{t \geq 1}$ are independent fair signs, each equal to $+1$ or -1 with probability one half, also independent of the trader's private random seed. If w_t is measurable with respect to the information available when the position is chosen and $w_t(P_{t+1} - P_t)$ is integrable, then

$$G_n = \sum_{t=0}^{n-1} w_t(P_{t+1} - P_t), \quad n = 0, \dots, T,$$

is a martingale: its expected next value, conditional on current information, equals its current value. In particular, $\mathbb{E}[G_T] = 0$, where $\mathbb{E}$ denotes expectation.

Proof. Let R denote the trader's private seed and set $\mathcal{G}_t = \sigma(P_0, \dots, P_t, R)$, the information generated by these observations and the seed; here $\sigma(\cdot)$ denotes their generated sigma-algebra, the collection of events they determine. Allowing $\mathcal{G}_t$ to contain all of R only strengthens the claim. Predictability makes w_t $\mathcal{G}_t$ -measurable, meaning determined by this information, while ξ_{t+1} is independent of $\mathcal{G}_t$ and has mean zero. Hence

$$\mathbb{E}[w_t(P_{t+1} - P_t) \mid \mathcal{G}_t] = \epsilon P_t w_t \mathbb{E}[\xi_{t+1} \mid \mathcal{G}_t] = 0.$$

Here $\mathbb{E}[\cdot \mid \mathcal{G}_t]$ is expectation conditional on that information. The integrability assumption therefore makes $(G_n)_{n=0}^T$ a martingale, and $\mathbb{E}[G_T] = \mathbb{E}[G_0] = 0$. $\square$

Private randomization prevents a deterministic countermarket based only on the public history from mechanically opposing the unseen current draw, but it does not yield strictly positive expected gain in every environment. More precisely, under the proposition’s integrability assumptions it excludes $\forall P : \mathbb{P}_R(G_T(P, R) > 0) = 1$ at the fixed horizon T , where $\mathbb{P}_R$ takes probability over the trader’s seed for a fixed path. Each of the finitely many fair-sign prefixes has positive probability and a computable infinite continuation; such a guarantee would imply $G_T > 0$ almost surely jointly, hence $\mathbb{E}[G_T] > 0$, a contradiction. It does not exclude a high win probability, profit under another law, or benchmark-relative success.

b. Optional online-game reading The recursion in Theorem 1 also defines a finite perfect-information game, in which both players observe every previous move: the trader announces w_t , the environment chooses the next return, and Corollary 4 gives the trader-first maximin value zero. This reading is optional because the fixed deterministic trader’s countermarket program already generates the same path offline. Borel determinacy concerns infinite games whose winning sets are generated from finite-history events by complements and countable unions (Borel sets); it does not by itself identify a winner here [23]. Hidden private randomization is a different information structure: Proposition 5 proves zero expected gain under one passive computable fair-sign law, not a value for a game with randomized strategies without specified strategy spaces.

E. A Passive Market: Martin-Löf Randomness

a. When does process randomness yield a random record? Process randomness describes how individual bits are generated and depend on their past. Martin-Löf (ML) randomness is a property of the resulting *infinite sequence* relative to a specified probability law. *Independent fair bits produce a fair-coin ML-random sequence with probability one.* Equivalently, the sufficient process condition is

$$\mathbb{P}(X_{n+1} = 1 \mid X_1, \dots, X_n) = \frac{1}{2} \quad \text{a.s. for every } n \geq 0, \quad (8)$$

Here X_n is the random bit at date n , $\mathbb{P}$ denotes probability, and “a.s.” means almost surely, or with probability one. The condition includes the first bit at $n = 0$.

Independence alone is insufficient for *fair-coin* randomness: independent bits with a fixed probability $p \neq 1/2$ of being one almost surely have limiting frequency p . Fair one-bit probabilities alone are also insufficient: repeating one fair bit forever gives fair marginals—each bit separately is equally likely to be zero or one—but a constant, predictable record. Independence is not necessary for ML randomness relative to the process’s *own* computable law, which may include bias or dependence.

For FAPP use, the conclusion is therefore conditional on a justified source model, for example independent fair

trials. Tests for bias and serial dependence (statistical dependence between different dates) can challenge that model over observed scales; passing them cannot certify independence indefinitely or the ML randomness of an infinite output. An almost-sure conclusion allows exceptional records of probability zero. A fixed deterministic pseudorandom generator with a finite seed cannot produce a fair-coin ML-random infinite sequence, however well its finite outputs test.

Here is the precise classical result. A computable law μ has cylinder probabilities $\mu([s])$ uniformly approximable to any requested accuracy, where s is a finite bit string and its cylinder $[s]$ comprises all infinite sequences beginning with s . An ML test is a uniformly computable enumerable sequence of open sets (U_k) with $\mu(U_k) \leq 2^{-k}$: each set is a union of cylinders, and one algorithm lists their defining prefixes given k . A record is μ -ML-random if it avoids $\bigcap_k U_k$, the set common to every level, for every such test. Write MLR_μ for the set of these random records.

Proposition 6 (Classical Measure-One Result, Martin-Löf). *If a binary process has computable joint law μ , then*

$$\mathbb{P}((X_1, X_2, \dots) \in \text{MLR}_\mu) = 1. \quad (9)$$

This includes independent fair trials, independent bits with a fixed computable probability of one (a Bernoulli bias), and computable dependent laws [24, Sec. IV].

Proof. Each test’s failure set has measure zero. There are countably many effective tests, so their union still has measure zero. $\square$

The law must be specified: randomness relative to an atomic law, which gives positive probability to individual records, need not mean fresh uncertainty at each trial. For example, either constant record is random for the law that repeats one fair bit forever.

b. Borel normality is weaker A sequence is Borel normal in base b when every word, or contiguous block, of k digits has limiting frequency b^{-k} , for every $k \geq 1$; the digit alphabet is $0, \dots, b-1$. ML randomness for uniform base- b digits implies normality, but the converse fails. Champernowne’s decimal $0.123456789101112\dots$, formed by concatenating the positive integers, is normal in base ten yet computable, hence not ML-random for the uniform decimal law [25]. Normality is difficult to establish for a given infinite record: one must control all block lengths in the infinite limit. No finite frequency check proves it, since any finite prefix admits both normal and nonnormal continuations. A proof from a construction, as for Champernowne, can establish normality; extensive sampling alone cannot.

F. Ramsey Regularity and the Finite-Pattern Trap

A large market record can be pattern-rich even when it contains no exploitable law. Ramsey theory makes

this finite combinatorial point without assuming a probability model. For fixed positive integers c, m, q , there is a finite number $R_q(m; c)$ such that every assignment of one of c colors to each q -element subset of any set of size at least $R_q(m; c)$ has an m -element subset whose q -subsets all receive the same color. Thus every red-blue coloring of the pairs among six objects contains a monochromatic triangle, with all three pairs the same color. Van der Waerden’s theorem similarly forces an arithmetic progression—equally spaced indices—of any prescribed finite length, all of one color, in every sufficiently long finite coloring [26, 27]. These thresholds may be enormous, but they are finite and computable, unlike a Busy-Beaver horizon.

The guarantee depends on the coding and is weaker than a trading signal. In a discretized return record, equal symbols at equally spaced dates need not be consecutive, form a price trend, have a favorable sign, or persist beyond the sample. Ramsey existence therefore asserts neither statistical dependence nor causation, predictability, or profit. Nor does it conflict with algorithmic randomness. A fair-coin Martin-Löf-random sequence is Borel normal, so every finite binary word, including arbitrarily long runs and finite chart motifs, occurs with its fair-coin limiting frequency. Randomness here means escaping every effective null test, not local patternlessness [28, 29].

Calude and Longo call regularities forced by data volume *Ramsey-type correlations*. Their incompressibility argument emphasizes that almost all sufficiently long finite strings resist compression by any fixed positive fraction while still containing the finite regularities forced by Ramsey theory [30]. This is an epistemic warning, not a theorem that every sample correlation is false: a Ramsey configuration is not itself a correlation coefficient, and a pattern observed in one finite history is compatible with both random and computable continuations.

Trading research compounds this pattern abundance with selection. A null hypothesis is the baseline claim assessed by a statistical test, such as the absence of a forecasting relation. If all M null hypotheses are true, their rejection events are independent, and each has rejection probability exactly α , the probability of at least one false rejection is $1 - (1 - \alpha)^M$; if their probabilities are at most α , this is an upper bound. Without independence, the union bound, which bounds the probability of any rejection by the sum of individual rejection probabilities, gives $\min\{1, M\alpha\}$ as an upper bound. This probabilistic multiple-testing problem is distinct from Ramsey’s deterministic existence theorem, but its practical lesson is aligned: a trend or cross-market relation retained after searching many assets, windows, transformations, lags, and parameters is a candidate hypothesis, not yet an edge. A defensible FAPP claim must report the search universe and survive selection correction (accounting for the search over alternatives), costs, an untouched hold-out (data never used to choose or tune the strategy) or independent replication, and testing outside the selected

regime [31].

G. What Rice’s Theorem Does and Does Not Say

Rice’s theorem makes every nontrivial extensional property of arbitrary partial computable functions undecidable [32]. “Extensional” means depending only on the function computed, not its program text; “nontrivial” means holding for some such functions but not all. In particular, totality is undecidable in general [12]. It does not make a fixed total strategy’s gain on a completed finite computable path undecidable, because that program can simply be simulated. Theorems 2 and 3 instead give the required promise-domain results by explicit Halting reductions: the unbounded questions of eventual profit and future-event occurrence are undecidable even though every code produced by the reductions is total.

III. THE NO-ARBITRAGE FOUNDATION

Here the market is no longer an opponent choosing a response after seeing the trader’s action. Instead, financial admissibility and no-arbitrage delimit which trader-market pairs are coherent.

A. Discounted Gains and Admissibility

Let the market be defined on a filtered probability space $(\Omega, \mathcal{F}, \{\mathcal{F}_t\}_{t \in [0, T]}, \mathbb{P})$. Here Ω is the set of possible outcomes, $\mathcal{F}$ the collection of events assigned probabilities by the physical, or real-world, law $\mathbb{P}$, and $\mathcal{F}_t$ the information available by time t . The increasing family $\{\mathcal{F}_t\}$ is the *filtration*, and T is the finite trading horizon. Let $B_t > 0$ be a numeraire and S_t the vector of traded asset prices. Write $\bar{S}_t = S_t/B_t$ for discounted prices. A predictable strategy chooses holdings from information available just before each price move; $\bar{S}$ -integrability means that its stochastic gain integral exists, which differs from finite expected absolute value. Such a strategy H has discounted self-financing gain

$$G_t = (H \cdot \bar{S})_t, \quad G_0 = 0. \quad (10)$$

Here H_t records the asset holdings, and the dot denotes stochastic integration: the continuous-time counterpart of summing holdings times discounted price changes. The strategy is *admissible* if there is a finite constant $a \geq 0$, a fixed bound on its possible intermediate loss, such that

$$G_t \geq -a \quad \text{for all } t \in [0, T], \quad \mathbb{P}\text{-a.s.} \quad (11)$$

This uniform lower bound rules out doubling systems, which repeatedly increase the stake after losses and can require unbounded credit.

Observation 7 (Universality Implies Classical Arbitrage). *If an admissible self-financing strategy H with zero initial capital satisfies*

$$G_T \geq 0 \quad \mathbb{P}\text{-a.s.}, \quad \mathbb{P}(G_T > 0) > 0,$$

then H is a classical arbitrage. In particular, a pointwise universal strategy is a classical arbitrage whenever its gain is defined on an admissible path class carrying $\mathbb{P}$.

Proof. The displayed payoff conditions, together with zero initial cost, self-financing, and admissibility, are the standard classical-arbitrage conditions. Pointwise strict positivity on every admissible path implies them, with $\mathbb{P}(G_T > 0) = 1$. $\square$

Remark 2 (Terminology and Strength). The pathwise condition is sometimes called a sure or strong arbitrage, but it is not synonymous with *arbitrage of the first kind*. The latter means that one fixed nonnegative terminal payoff, positive with positive probability, can be covered from every strictly positive initial capital while wealth remains nonnegative. Its absence, No Arbitrage of the First Kind (NA1), is equivalent to No Unbounded Profit with Bounded Risk (NUPBR) under the standard portfolio hypotheses [33]. Classical no-arbitrage (NA) alone already excludes the universal strategy in Observation 7; No Free Lunch with Vanishing Risk (NFLVR), explained next, supplies the standard martingale-measure formulation.

B. NA, NUPBR, and Benchmark-Relative Gains

A *stopping time* is a decision time recognizable from the information then available. A *local martingale* becomes a martingale after stopping along a suitable increasing sequence of such times that eventually covers the horizon. A *semimartingale* is the sum of a local martingale and a process of finite variation, meaning finite accumulated absolute movement on each finite interval; this is the standard framework for stochastic trading gains.

The relevant no-arbitrage notions do not form the simple chain sometimes suggested in informal discussions. NUPBR requires terminal gains from zero-cost strategies whose gains never fall below minus one to be *bounded in probability*: the probability of exceeding a large gain threshold must tend to zero uniformly over those strategies. NFLVR excludes even approximate arbitrage: bounded payoffs covered by admissible zero-cost terminal gains cannot approach a nonnegative payoff, positive with positive probability, with the largest possible approximation error tending to zero apart from probability-zero exceptions. In the standard semimartingale setting, NFLVR combines classical NA with NUPBR, whereas NA and NUPBR are in general logically distinct. NUPBR is equivalent, under standard portfolio hypotheses, to the existence of a *numeraire portfolio*: when other positive wealth processes are divided

by its wealth, the resulting relative-wealth processes are *supermartingales*, whose conditional expected future values do not exceed their present values [33, 34]. It may hold even when the stronger NFLVR condition and the equivalent local martingale measure defined below fail.

This distinction prevents two overstatements. First, Observation 7 uses only classical NA; NUPBR alone is not silently substituted for it. If a zero-cost gain were nonnegative at all intermediate times and positive at terminal time, scaling would violate NUPBR, but terminal pathwise positivity together with only a strategy-specific lower bound does not by itself supply that scaling argument. Second, benchmark and stochastic-portfolio approaches can generate gains *relative* to a numeraire or market portfolio without producing zero-cost positive gain on every path. The market portfolio weights stocks by their shares of total market value; a diverse equity market keeps every such weight uniformly below one. Relative arbitrage matches a chosen benchmark at terminal time with probability one and beats it with positive probability, starting from the same wealth. Real-world benchmark pricing values payoffs in benchmark units using the physical probability law. These approaches therefore do not contradict the theorem [35, 36].

C. The Martingale Constraint

A price process is *locally bounded* if it is bounded by a deterministic constant after stopping along an increasing sequence of stopping times that eventually covers the horizon. For a locally bounded semimartingale price process, the Delbaen–Schachermayer form of the First Fundamental Theorem of Asset Pricing (FTAP) states that NFLVR is equivalent to the existence of an equivalent local martingale measure (ELMM) $\mathbb{Q} \sim \mathbb{P}$ under which $\bar{S}$ is a local martingale [37, 38]. Here $\mathbb{Q}$ is a pricing probability law, and equivalence, denoted by $\sim$, means that the two laws assign probability zero to exactly the same events. For general unbounded semimartingales the corresponding formulation uses an equivalent σ -martingale measure: under that law, some strictly positive predictable weighting of price increments has a martingale stochastic integral [39]. The corollary below retains the locally bounded setting.

Corollary 8 (No Universal Zero-Cost Strategy). *If an ELMM $\mathbb{Q}$ exists, no admissible self-financing strategy with zero initial capital can satisfy $G_T \geq 0$, $\mathbb{P}$ -almost surely, and $\mathbb{P}(G_T > 0) > 0$. Consequently, no pointwise universal strategy exists among admissible strategies in this stochastic model.*

Proof. Under $\mathbb{Q}$, the admissible stochastic integral $G = H \cdot \bar{S}$ is a local martingale bounded from below and therefore a supermartingale. Hence $\mathbb{E}_{\mathbb{Q}}[G_T] \leq G_0 = 0$. If the displayed arbitrage conditions held, equivalence would give $G_T \geq 0$, $\mathbb{Q}$ -almost surely, with $\mathbb{Q}(G_T > 0) > 0$. Consequently $\mathbb{E}_{\mathbb{Q}}[G_T] > 0$, a contradiction. $\square$

The ELMM conclusion concerns discounted gains under the pricing measure $\mathbb{Q}$. It does not say that every risky strategy has non-positive expected return under the physical measure $\mathbb{P}$; positive physical expected returns may compensate risk.

D. Win Rate Is Not Expected Value

The probability of closing a trade at a profit can greatly exceed $1/2$ even in a fair market. Optional stopping—preservation of a martingale’s expected value under suitable stopping conditions—gives the precise distinction.

Proposition 9 (Asymmetric Barriers in a Fair Random Walk). *Let X_n be a simple symmetric random walk with $X_0 = 0$: each step is independently $+1$ or -1 , with equal probabilities. For positive integers a, b , let τ be the first time either barrier is reached:*

$$\tau = \inf\{n \geq 0 : X_n = a \text{ or } X_n = -b\}.$$

Then

$$\mathbb{P}(X_\tau = a) = \frac{b}{a+b}, \quad \mathbb{E}[X_\tau] = 0. \quad (12)$$

Proof. For this finite interval, $\tau < \infty$ almost surely and $|X_{n \wedge \tau}| \leq \max\{a, b\}$, where $n \wedge \tau$ is the earlier of n and τ . Optional stopping gives $\mathbb{E}[X_{n \wedge \tau}] = 0$, and bounded convergence—which permits exchanging the limit and expectation for these uniformly bounded variables—yields $\mathbb{E}[X_\tau] = 0$ [40]. If $p = \mathbb{P}(X_\tau = a)$, then $0 = \mathbb{E}[X_\tau] = ap - b(1 - p)$, which gives $p = b/(a + b)$. $\square$

With a take-profit order closing the trade at gain $a = 1$ and a stop-loss order closing it at loss $b = 9$, the win probability is 0.9, but the expected payoff is zero: nine frequent unit gains are offset by one nine-unit loss in expectation. Thus prediction accuracy, trade win rate, and expected discounted profit are three different quantities.

IV. THE COMBINATORIAL PERSPECTIVE

Here the opposing side is an ensemble, a collection of possible market records equipped with a probability law, rather than an adaptive market. The optimization NFL theorem and the elementary prediction result needed for market-direction claims should be distinguished.

Theorem 10 (Classical Optimization NFL, Wolpert—Macready). *Let X be a finite search domain and Y a finite set of possible function values. Give Y^X , the class of all functions from X to Y , its uniform law, so each function is equally likely. A non-revisiting black-box algorithm learns function values only by querying inputs and never queries the same input twice. For any two such algorithms and any fixed number $q \leq |X|$ of queries, where $|X|$ is the number of inputs, the resulting q -term*

sequences of observed values have the same distribution. Consequently, every performance functional, meaning a numerical score depending only on that observed value sequence, has the same average for both algorithms [41].

The theorem concerns optimization using an oracle, a device that returns a function’s value at each queried input. It does not, by itself, imply a statement about self-financing gains on stochastic price paths. A sharpening makes its structural assumption explicit: for a uniform distribution on a finite function class, the NFL equality for all non-revisiting algorithms holds precisely when the class is closed under permutations of the search domain: relabeling the inputs of any function leaves it within the same class [42]. Financial path classes are generally not permutation-closed: temporal order, continuity, financing, and no-arbitrage restrictions distinguish one permutation from another. Thus a financial edge, when present, resides in exactly such symmetry-breaking structure and the inductive bias used to model it, meaning the assumptions that guide generalization beyond observed data.

For binary directional prediction the relevant result is simpler and can be derived directly.

Proposition 11 (Uniform Binary Prediction). *Let $X_1, \dots, X_T$ be uniformly distributed on $\{0, 1\}^T$, so every length- T binary string is equally likely. A causal deterministic predictor uses only past observations to choose its forecast $\hat{X}_t = f_t(X_1, \dots, X_{t-1})$. Its mean accuracy*

$$A_T = \frac{1}{T} \sum_{t=1}^T \mathbf{1}\{\hat{X}_t = X_t\}$$

uses the indicator $\mathbf{1}$, equal to one for a correct forecast and zero otherwise, and satisfies $\mathbb{E}[A_T] = 1/2$. The same holds for a randomized predictor driven by a private random seed independent of $(X_1, \dots, X_T)$.

Proof. Conditional on every history $X_1, \dots, X_{t-1}$, the next bit X_t is an independent fair bit. Therefore $\mathbb{P}(\hat{X}_t = X_t \mid X_1, \dots, X_{t-1}) = 1/2$. Linearity of expectation gives the result. In the randomized case, conditioning additionally on the independent private seed leaves X_t fair given the history and proves the same equality. $\square$

Consequently, no predictor can achieve expected directional accuracy greater than $1/2$ under the uniform distribution over all binary sequences; in particular, none can exceed $1/2$ on every sequence. This is an NFL symmetry statement, not a model of actual markets. Financial sign sequences are not uniform in general, and a directional forecast does not determine trading profit because payoff sizes, financing, and stopping rules matter (Proposition 9).

The defensible practical conclusion is conditional: a claim of repeatable success must identify non-uniform structure in the physical measure $\mathbb{P}$, such as a trend, mean reversion, a risk premium, information, or an institutional constraint that breaks the NFL symmetry. The hypothesized edge concerns a restricted environment, not universal algorithmic superiority.

V. PRACTICAL ROBUSTNESS DIAGNOSTICS: TIME REVERSAL AND PERMUTATIONS

This section turns the trader–market leitmotif into a falsification tool. Hold the trader fixed and transform only the market record. Among the countable family of computably specified stress transformations, time reversal is one of the simplest, particularly for detecting dependence on directional persistence. It is not a further impossibility theorem on the level of the Halting reductions or no-arbitrage.

Let $\mathcal{M}$ be a space of strictly positive discounted price paths $P: [0, T] \rightarrow \mathbb{R}_{>0}^n$, where $\mathbb{R}_{>0}^n$ denotes vectors of n positive real prices, one per asset, and T is the horizon. The time reversal of P is

$$\tilde{P}(t) = P(T - t), \quad t \in [0, T]. \quad (13)$$

Literal reversal reuses the observed price levels and preserves the magnitudes of log-price moves, the logarithms of successive price ratios, while reversing their order and signs. It is not the same as merely negating arithmetic returns: a gross return R is the later price divided by the earlier price, and the corresponding arithmetic return is $R - 1$. The reversed gross return is $1/R$. In an empirical implementation the strategy must be rerun causally from scratch on the reversed record; reversing the original orders or realized profit would leak future information. If the initial level is prescribed, literal reversal can leave the market class. One may instead test componentwise normalized reversal,

$$P_i^{\text{rev}}(t) = P_i(0) \frac{P_i(T - t)}{P_i(T)},$$

where i labels the asset and rev its reversed path, provided it remains in the claimed class. Scaling can change the actions of a strategy that is not scale invariant, meaning that its decisions change when all prices of an asset are multiplied by the same positive constant; it is therefore a distinct test. For a causal strategy σ , let $\Pi_t[\sigma; P]$ denote its discounted gain from zero initial capital through time t on P . Call σ *pathwise admissible* on $\mathcal{M}$ if some finite C satisfies $\Pi_t[\sigma; P] \geq -C$ for every $P \in \mathcal{M}$ and $t \in [0, T]$.

Definition 1 (Pathwise Universal Strategy). A causal strategy σ is *pathwise universal* on $\mathcal{M}$ if it is self-financing and pathwise admissible there and satisfies $\Pi_T[\sigma; P] > 0$ for every $P \in \mathcal{M}$.

Criterion 1 (Reversal Falsification Test). Suppose $\mathcal{M}$ is closed under time reversal, meaning that reversing any member produces another member. If there exists $P \in \mathcal{M}$ such that

$$\min\{\Pi_T[\sigma; P], \Pi_T[\sigma; \tilde{P}]\} \leq 0, \quad (14)$$

then σ is not pathwise universal on $\mathcal{M}$.

Proof. Both P and $\tilde{P}$ belong to $\mathcal{M}$. A non-positive gain on either member of the pair contradicts Definition 1. $\square$

One non-winning path refutes a claim quantified over every path. Reversal supplies a paired candidate, not a proof that one member must fail: a causal strategy may take different positions on the reversed history. Diagonalization and no-arbitrage provide the genuine impossibility results under their respective assumptions.

a. Permuting a finite record Time reversal is a special permutation of a sampled record. For a finite record of discounted price levels $P = (P_0, \dots, P_N)$, let π be a permutation of $\{0, \dots, N\}$, a rearrangement using every index exactly once, and define

$$(P^\pi)_t = P_{\pi(t)}, \quad t = 0, \dots, N. \quad (15)$$

Thus $\pi(t) = N - t$ gives reversal. A fixed algorithm or a random shuffle can choose π . For vector-valued prices, each complete price vector moves together. Let $\mathcal{M}_N$ be the claimed class of such discrete records and $\Pi_N[\sigma; P]$ the discounted gain from zero initial capital when the same causal rule is rerun from its prescribed initial state on P . Use the discrete version of Definition 1 on this class.

Criterion 2 (Permutation Falsification Test). A pathwise universal strategy on $\mathcal{M}_N$ must satisfy

$$\begin{aligned} \Pi_N[\sigma; P^\pi] &> 0 \\ \text{whenever } P \in \mathcal{M}_N \text{ and } P^\pi \in \mathcal{M}_N. \end{aligned} \quad (16)$$

One permitted permutation with zero or negative gain refutes that universal claim. The criterion requires positivity, not equality of profits across permutations.

Proof. Every retained P^π is a member of the claimed market class, so strict profit on every member requires strict profit on this one. $\square$

The shuffle rearranges price levels, not returns or previously chosen orders. At each date the rerun uses only the history available in the new order. Permuting levels preserves their multiset, the collection of levels with repetitions counted, but generally changes the returns and their magnitudes. It can also violate a fixed initial-price condition, a bound on price moves, or another restriction defining $\mathcal{M}_N$. Only transformations that remain within that class test its universal claim; closure under all permutations is unnecessary. Failure after destroying a trend outside a strategy's stated domain does not refute a claim confined to that domain. Passing finitely many shuffles of one record does not establish universal profitability on a larger path class.

Unlike the permutation closure used in optimization NFL (Section IV), this diagnostic assumes no uniform distribution or equality of average performance. Randomness in the choice of stress records alone supplies neither a market law nor a statistical significance level.

Greg Egan's *Permutation City* (1994) [43] provides a literary analogy by exploring the temporal ordering

of simulated computational states. Here permutations change the trader's observation sequence; the causal rule is rerun on each record, and Criterion 2 follows from the universal-profit requirement.

Remark 3 (Computable Stress Tests Beyond Reversal). Time reversal is one of countably many computable probes: reversing return signs, joining segments from different market conditions (regime splicing), adding price jumps or execution delays, changing the size of price fluctuations (volatility rescaling), and varying financing or transaction costs. Since program totality is undecidable, an operational battery must select known-terminating transformations whose outputs remain in the claimed market class. Here “test” means empirical stress, not a Martin-Löf test: a uniformly computably enumerable sequence of open sets with effective measure bounds. Reversal defines no null set and certifies no randomness; passing it shows only tested robustness, not universality.

Remark 4 (Filtrations and Time Reversal). Stochastic model classes are not automatically reversal-closed: $P(T-t)$ is generally not adapted to the forward filtration, meaning that its value at time t need not be determined by information available by then. Reversing a diffusion, a continuous process driven by Brownian noise, requires additional hypotheses [44]. For example, if

$$X_t = \log(S_t/S_0) = mt + \sigma W_t,$$

X_t is the stock's cumulative log return, m its drift (average rate of change), and W_t standard Brownian motion, whose independent normal increments have mean zero and variance equal to elapsed time. Here $\sigma \geq 0$ denotes volatility, the scale of random fluctuations, rather than a trading strategy. Then $X_{T-t} - X_T$, under the information revealed by the reversed increments, has drift $-m$ and volatility σ . Thus a class containing both drift signs is closed in law under this transformation: the transformed process has a distribution within the class. Under the physical law, $m = \mu - \sigma^2/2$, where μ is the stock's instantaneous expected proportional return. Under a risk-neutral measure, a probability law used for pricing through discounted expectations, the discounted stock may be a martingale while its logarithm retains the Itô drift $-\sigma^2/2$, the correction caused by taking logarithms of Brownian-driven prices. Neither an ELMM nor zero drift in an additive Brownian price model establishes detailed balance, meaning invariance of equilibrium history probabilities under time reversal, in general. The diagnostic remains pathwise; Section III gives the probabilistic constraints.

VI. CASE STUDY: THE WHEEL OPTIONS STRATEGY

The Wheel makes the trader-market configuration concrete. A put gives its buyer the right to sell stock at a specified exercise price (the strike); a call gives the right

to buy. The Wheel first sells a cash-secured put, reserving enough cash to buy the stock if assigned, that is, required to fulfill the seller's obligation. It then sells a covered call, backed by the acquired stock. For one matched pair of European options, exercisable only at their expiry date (maturity) T , with strike K , the terminal-payoff identity is

$$S_T - (S_T - K)^+ = K - (K - S_T)^+ = \min\{S_T, K\}. \quad (17)$$

Here S_T is the terminal stock price and $x^+ = \max\{x, 0\}$ denotes the positive part of x . The cash leg pays K at maturity and must be financed at its present value, the amount invested today to deliver K then; dividends and other cash flows must be matched separately. This is a local payoff comparison, not a model of the rolling Wheel. Early exercise of American options, assignment, changing strikes and maturities, rules for replacing existing options (rolling), bid-ask spreads (the difference between quoted purchase and sale prices), taxes, and financing make the realized strategy depend on the intervening price history. Each sold, or short, option nevertheless carries the associated downside exposure.

Option selling can nevertheless have positive mean return under a physical law. The difference between risk-neutral and physical expectations of future variance is commonly called the variance risk premium, with sign conventions varying across the literature, and option-market evidence documents compensation for bearing volatility and jump risk [45, 46]. Such compensation is empirical and state-dependent: its sign and magnitude depend on option prices, strikes, transaction costs, and the physical distribution.

The local payoff comparison suggests four stress scenarios for the rolling strategy. A crash after assignment produces essentially one-for-one downside below K . A persistent decline can make repeated premiums too small to offset the renewed short-put exposure. In a large rally, assignment and call exercise can still leave a positive absolute profit, but the cap at K can underperform a buy-and-hold benchmark. Finally, under proportional position sizing, capital exhaustion is absorbing: if wealth evolves as $W_{t+1} = W_t(1 + fR_{t+1})$, where f is the fixed fraction exposed to the risky arithmetic return R_{t+1} , then $1 + fR_{t+1} = 0$ leaves wealth at zero thereafter unless external capital is injected. A fully cash-secured, unlevered implementation, using no borrowed investment capital, has bounded loss, but an underlying that falls to zero can consume almost all committed capital net of premium.

A rally/crash reversal is therefore a useful paired stress test, although the payoff identity, rather than reversal alone, supplies the economic mechanism. Repeated small premiums also need not imply satisfactory time-average growth:

$$\frac{1}{T} \log \frac{W_T}{W_0} = \frac{1}{T} \sum_{t=1}^T \log(1 + fR_t),$$

so frequent gains can coexist with poor long-run growth when rare losses are large [47]. The Wheel may be a

defensible FAPP strategy under a specified risk premium, capital policy, cost model, and benchmark; it is not a pathwise profit guarantee.

VII. UNIVERSAL PORTFOLIOS AND DISTANCE FROM UNIVERSALITY

Here the trader changes the claim made against the market side: it seeks small regret to a benchmark portfolio class, not positive profit on every path. The word “universal” also appears in a celebrated positive theorem. For m assets, let the price-relative vector $x_t \in \mathbb{R}_{>0}^m$ contain each asset’s price at date t divided by its price at date $t - 1$. A constant-rebalanced portfolio restores the same nonnegative investment fractions $b = (b_1, \dots, b_m)$, with $\sum_i b_i = 1$, at each date. Starting from unit wealth, it has wealth

$$S_T(b) = \prod_{t=1}^T b^\top x_t.$$

The transpose notation $b^\top x_t = \sum_i b_i x_{t,i}$ gives the portfolio’s gross return for that period. Cover’s universal portfolio is nonanticipating, using only previously observed prices, and, under the conditions of his theorem, has the same asymptotic exponential growth rate, measured by log wealth per period, as the best constant-rebalanced portfolio chosen in hindsight on every individual bounded sequence [48]. Writing S_T^{UP} for its wealth and calling the benchmark’s excess log wealth the log regret,

$$\frac{1}{T} \left[\log \max_b S_T(b) - \log S_T^{\text{UP}} \right] \longrightarrow 0. \quad (18)$$

Equation (18) is a relative guarantee, not an absolute-profit guarantee. If every constant-rebalanced portfolio loses money on a path, the universal portfolio may lose as well while retaining vanishing per-period regret. Cover’s result therefore does not contradict Corollary 8; it illustrates a useful benchmark-relative meaning of “universal” once the comparison class is specified.

Online regret methods, which update decisions as observations arrive, also yield financial guarantees of another kind. DeMarzo, Kremer, and Mansour show that gradient strategies, adjusting decisions using local rates of change of the objective, can minimize asymptotic regret and imply trading strategies that provide arbitrage-based option price bounds. They require neither continuous price paths, complete markets (in which every contingent payoff can be replicated by trading), nor a pricing kernel (a random factor converting future payoffs into present prices through expectation). The bounds depend on realized quadratic variation, the limiting sum of squared price changes as observation intervals shrink [49]. This is robust pricing and hedging, not guaranteed positive profit or next-move prediction.

a. A Hierarchy of Attainable Claims The useful alternatives to pathwise universality are not weaker points on one single order; they change the benchmark, horizon, or probability quantifier. Table II records the distinctions.

Superhedging gives another quantitative notion: once a contingent claim, a payoff depending on future market outcomes, and an admissible model class are fixed, its superhedging price is the least initial capital needed to ensure that terminal wealth covers that payoff on every allowed path. Requiring zero initial capital while demanding a strictly positive payoff is precisely the degenerate, arbitrage-producing endpoint ruled out above.

VIII. FAPP STRATEGIES AND THEIR HONEST SCOPE

Although pathwise universal profit is excluded, a strategy may be profitable *for all practical purposes* (FAPP, following Bell [50]) relative to a physical law $\mathbb{P}$, a benchmark, and a finite horizon. Market making, posting buy and sell quotes, may earn spreads when adverse selection (losses from trading with better-informed counterparties) and inventory risk (price risk on accumulated holdings) are controlled. Kelly investing chooses position sizes to maximize expected logarithmic wealth growth when the relevant distribution is sufficiently well specified [51]. Neither is an unconditional guarantee.

Here “insight” means a falsifiable restriction on the market side. In actual trading it usually appears as a heuristic, a useful but fallible decision rule: persistence, mean reversion, valuation correction (prices moving toward an estimated fundamental value), or compensation for an identified risk. The fact that a heuristic is algorithmic does not disqualify it; its edge is a regime-dependent empirical hypothesis. A repeatable claim stated before outcomes are observed must therefore specify its market class, horizon, costs, benchmark, and validation protocol.

Only lawfully acquired and used information is included in this practical strategy class; an informational advantage remains short of a pathwise guarantee.

The unit of performance matters as well. For a funded portfolio, let R^{nom} be the proportional total return before adjusting for inflation, R^{real} the return adjusted for purchasing power, and π inflation over the same interval under a stated price index. Then

$$1 + R^{\text{real}} = \frac{1 + R^{\text{nom}}}{1 + \pi}. \quad (19)$$

A positive nominal return can therefore be a purchasing-power loss. The price index is an economic deflator, converting currency values into purchasing-power units, not necessarily a tradable numeraire for the FTAP; the consumption benchmark and financing account serve different purposes.

Productivity growth may support positive real returns when its benefits reach investors as cash flows [52]. Prices,

TABLE II. Different meanings of a successful or universal strategy.

Notion	Guarantee	Status
Pathwise universal profit	Positive discounted zero-cost gain on every admissible path	Excluded among admissible strategies in the specified stochastic model under classical NA
Relative arbitrage	Outperformance of a specified market or numeraire portfolio	Possible under structural market assumptions
Cover universality	Vanishing per-period log regret to the best constant-rebalanced portfolio	Pathwise but benchmark-relative
FAPP or statistical edge	Positive performance under a specified physical law, data protocol, and horizon	Conditional and empirically falsifiable

discount rates, and the displacement of incumbents intervene [53]; aggregate growth therefore establishes neither next-move prediction nor active alpha, excess return attributable to the strategy relative to a specified risk-adjusted benchmark.

Equilibrium theory likewise permits only conditional edges. Grossman and Stiglitz show that perfectly informationally efficient markets are inconsistent with costly information acquisition: if prices revealed all information without reward, agents would not pay to produce it [54]. Conversely, the Milgrom–Stokey no-trade theorem limits trade based on private information alone [55]. Its assumptions include a common prior (shared initial probabilistic beliefs), common knowledge (each agent knows the relevant facts, knows that the others know, and so on), risk aversion (a preference for less risk at a given expected payoff), and Pareto optimality (an allocation that cannot improve one agent’s welfare without harming another). A complementary market-selection result is also conditional: among agents in Sandroni’s complete-market model who share an intertemporal discount factor, the weight attached to future utility (the model’s measure of satisfaction), and choose saving within the model, accurate forecasters prosper and inaccurate ones are driven out [56]. These equilibrium results depend on their model assumptions; none is a pathwise trading guarantee or a substitute for the FTAP.

Deployment can also change the market side. A large, crowded, or public strategy consumes liquidity, the capacity to trade near prevailing prices, moves quotes, attracts imitation or opposition, and changes the distribution on which it was estimated. Lo’s Adaptive Markets Hypothesis provides this evolutionary interpretation [57]: competition and impact may erode an edge without producing the tailored countermarket of Theorem 1. Feedback invalidates extrapolation that assumes unchanged performance as trading size grows; it does not imply that every large strategy loses.

Finally, the finite-pattern trap of Section II F becomes operational in strategy discovery. Trying many specifications and retaining the best backtest, a simulation on historical data, turns its Sharpe ratio into an order statistic, a value selected by its rank among the tested results.

The Sharpe ratio is mean return above the financing rate divided by the standard deviation of returns, which measures their spread around the mean. The Deflated Sharpe Ratio adjusts the assessment of significance for selection, departures from a normal return distribution, and the number of trials [31]; untouched validation data, costs, and deployment analysis remain necessary. These controls do not prove an edge, but make a finite-sample FAPP claim more honestly falsifiable.

IX. CONCLUSION

Everything returns to the same pair: a trader and a market. For every fixed total deterministic computable trader, Theorem 1 constructs a fixed positive computable price path whose next bounded proportional move opposes every nonzero position and remains nonzero during inactivity. The quantifiers are $\forall\sigma \exists\mathcal{E}_\sigma$: no physical market must observe the deployed orders, the always-flat trader earns zero, and every nonzero position loses on its trader-relative countermarket. This direct diagonal result uses no Halting-Problem reduction and identifies no single path that defeats every trader.

The remaining limits concern different market sides. Gold precludes identifying every computable binary generator in the limit, while next-bit diagonalization defeats each fixed computable predictor on its predictor-relative record. Halting reductions preclude uniform certification of eventual profit or encoded future events, and Busy Beaver excludes a computable description-size waiting horizon despite finitely solvable low-complexity cases. Under a stated computable reference measure, the generated record is Martin-Löf random almost surely; neither this property nor Borel normality can be certified from a finite sample. Ramsey regularity explains why finite structured subsequences nevertheless remain abundant, including in random records. No-arbitrage and uniform binary averaging supply still different boundaries. Their conclusions are complementary because their inputs, quantifiers, and success criteria differ.

Positive results arise only after restricting the claim. The Wheel may earn a risk premium while retaining put-

like downside; Cover guarantees asymptotic performance relative to a benchmark class; and FAPP strategies may exploit a specified regime, lawful information, or compensated risk. Such performance must be evaluated after costs and financing, against an explicit benchmark and, where relevant, in purchasing-power terms. Productivity-supported real growth is one possible regime hypothesis, not a next-move predictor or a guarantee of active alpha.

Any repeatable-success claim must therefore state its market restriction, probability law, benchmark, admissibility and financing conditions, horizon, data protocol, and deployment effects. Time reversal and other within-class permutations are robustness probes: one nonpositive gain refutes an all-path claim, while passing a finite set of probes does not certify universality. Useful strategies may exist, but no total deterministic computable self-financing rule guarantees strict gain on every positive computable price path; under an arbitrage-free stochastic model, no admissible zero-cost rule delivers nonnegative discounted gain almost surely and positive gain with positive probability.

ACKNOWLEDGMENTS

This research was funded in whole or in part by the Austrian Science Fund (FWF) [Grant DOI: 10.55776/PIN5424624]. For open access purposes, the author has applied a CC BY public copyright license to any author accepted manuscript version arising from this submission. The author acknowledges TU Wien Bibliothek for financial support through its Open Access Funding Programme.

OpenAI Codex was used to assist with literature organization, LaTeX editing, and checks of derivations. The author specified the scientific direction and prompts; model output retained in the manuscript was checked against the cited sources and explicit calculations. The author accepts full responsibility for the final text.

DATA AVAILABILITY

This is a purely mathematical work, and no data or software were created or analyzed in this study.

- [1] E. M. Gold, Limiting recursion, *The Journal of Symbolic Logic* **30**, 28 (1965).
- [2] E. M. Gold, Language identification in the limit, *Information and Control* **10**, 447 (1967).
- [3] L. Bachelier, Théorie de la spéculation, *Annales Scientifiques de l'École Normale Supérieure* **17**, 21 (1900).
- [4] P. A. Samuelson, Proof that properly anticipated prices fluctuate randomly, *Industrial Management Review* **6**, 41 (1965).
- [5] E. F. Fama, Efficient capital markets: A review of theory and empirical work, *The Journal of Finance* **25**, 383 (1970).
- [6] S. Legg, *Is there an Elegant Universal Theory of Prediction?*, Tech. Rep. IDSIA-12-06 (Istituto Dalle Molle di Studi sull'Intelligenza Artificiale, 2006) [arXiv:cs/0606070 \[cs.AI\]](https://arxiv.org/abs/cs/0606070).
- [7] R. J. Solomonoff, A formal theory of inductive inference. part I, *Information and Control* **7**, 1 (1964).
- [8] R. J. Solomonoff, A formal theory of inductive inference. part II, *Information and Control* **7**, 224 (1964).
- [9] R. J. Solomonoff, Complexity-based induction systems: Comparisons and convergence theorems, *IEEE Transactions on Information Theory* **24**, 422 (1978).
- [10] M. Hutter, Convergence and error bounds for universal prediction of nonbinary sequences, in *Machine Learning: ECML 2001*, Lecture Notes in Computer Science, Vol. 2167, edited by L. De Raedt and P. Flach (Springer, 2001) pp. 239–250.
- [11] V. V. V'yugin, On impossibility of sequential algorithmic forecasting, in *Kolmogorov Complexity and Applications*, Dagstuhl Seminar Proceedings, Vol. 6051, edited by M. Hutter, W. Merkle, and P. M. B. Vitányi (Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl, Germany, 2006) pp. 1–7.
- [12] A. M. Turing, On computable numbers, with an application to the Entscheidungsproblem, *Proceedings of the London Mathematical Society* **2**, 42, 230 (1937).
- [13] L. Blum and M. Blum, Toward a mathematical theory of inductive inference, *Information and Control* **28**, 125 (1975).
- [14] D. Angluin and C. H. Smith, Inductive inference: Theory and methods, *ACM Computing Surveys* **15**, 237 (1983).
- [15] L. M. Adleman and M. Blum, Inductive inference and unsolvability, *The Journal of Symbolic Logic* **56**, 891 (1991).
- [16] A. Svozil and K. Svozil, Induction and physical theory formation as well as universal computation by machine learning, in *Swarm Intelligence: From Social Bacteria to Humans*, edited by A. Schumann (CRC Press, Taylor & Francis Group, LLC, Boca Raton, FL, 2021) pp. 170–180, [arXiv:1609.03862](https://arxiv.org/abs/1609.03862).
- [17] C. D. Moore, Unpredictability and undecidability in dynamical systems, *Physical Review Letters* **64**, 2354 (1990), cf. Ch. Bennett, *Nature*, **346**, 606 (1990).
- [18] T. Rado, On non-computable functions, *The Bell System Technical Journal* **XLI(41)**, 877 (1962).
- [19] G. J. Chaitin, Computing the busy beaver function, in *Open Problems in Communication and Computation*, edited by T. M. Cover and B. Gopinath (Springer, New York, 1987) p. 108.
- [20] C. S. Calude, M. J. Dinneen, and C.-K. Shu, Computing a glimpse of randomness, *Experimental Mathematics* **11**, 361 (2002), [arXiv:nlin/0112022](https://arxiv.org/abs/nlin/0112022).
- [21] M. Barnsley, *Fractals Everywhere* (Academic Press, San Diego, 1988).
- [22] C. S. Calude, E. Calude, and K. Svozil, The complexity of proving chaoticity and the Church–Turing thesis, *Chaos: An Interdisciplinary Journal of Nonlinear Science* **20**, 037103 (2010).
- [23] D. A. Martin, Borel determinacy, *Annals of Mathematics* **102**, 363 (1975).
- [24] P. Martin-Löf, The definition of random sequences, *Information and Control* **9**, 602 (1966).
- [25] D. G. Champernowne, The construction of decimals normal in the scale of ten, *Journal of the London Mathematical Society* **s1-8**, 254 (1933).
- [26] R. L. Graham, B. L. Rothschild, and J. H. Spencer, *Ramsey Theory*, 2nd ed. (John Wiley & Sons, New York, 1990).
- [27] R. Graham and J. H. Spencer, Ramsey theory, *Scientific American* **263**, 112 (1990).
- [28] C. S. Calude, *Information and Randomness—An Algorithmic Perspective*, 2nd ed. (Springer, Berlin, 2002).
- [29] A. Nies, *Computability and Randomness*, Oxford Logic Guides, Vol. 51 (Oxford University Press, Oxford, 2009).
- [30] C. S. Calude and G. Longo, The deluge of spurious correlations in big data, *Foundations of Science* **22**, 595 (2017), first published online 7 March 2016, [CDMTCS-488](https://doi.org/10.1007/s11464-016-0488-8).
- [31] D. H. Bailey and M. López de Prado, The deflated sharpe ratio: Correcting for selection bias, backtest overfitting, and non-normality, *The Journal of Portfolio Management* **40**, 94 (2014).
- [32] H. G. Rice, Classes of recursively enumerable sets and their decision problems, *Transactions of the American Mathematical Society* **74**, 358 (1953).
- [33] C. Kardaras, Market viability via absence of arbitrage of the first kind, *Finance and Stochastics* **16**, 651 (2012).
- [34] I. Karatzas and C. Kardaras, The numéraire portfolio in semimartingale financial models, *Finance and Stochastics* **11**, 447 (2007).
- [35] R. Fernholz, I. Karatzas, and C. Kardaras, Diversity and relative arbitrage in equity markets, *Finance and Stochastics* **9**, 1 (2005).
- [36] E. Platen and D. Heath, *A Benchmark Approach to Quantitative Finance*, Springer Finance (Springer, Berlin, 2010).
- [37] J. M. Harrison and S. R. Pliska, Martingales and stochastic integrals in the theory of continuous trading, *Stochastic Processes and their Applications* **11**, 215 (1981).
- [38] F. Delbaen and W. Schachermayer, A general version of the fundamental theorem of asset pricing, *Mathematische Annalen* **300**, 463 (1994).
- [39] F. Delbaen and W. Schachermayer, The fundamental theorem of asset pricing for unbounded stochastic processes, *Mathematische Annalen* **312**, 215 (1998).
- [40] J. L. Doob, *Stochastic Processes* (John Wiley & Sons, New York, 1953).
- [41] D. H. Wolpert and W. G. Macready, No free lunch theorems for optimization, *IEEE Transactions on Evolutionary Computation* **1**, 67 (1997).
- [42] C. Schumacher, M. D. Vose, and L. D. Whitley, The no free lunch and problem description length, in *Proceedings of the Genetic and Evolutionary Computation Conference*

- (Morgan Kaufmann, San Francisco, 2001) pp. 565–570.
- [43] G. Egan, *Permutation City* (Orion/Millennium, London, 1994).
 - [44] U. G. Haussmann and É. Pardoux, Time reversal of diffusions, *The Annals of Probability* **14**, 1188 (1986).
 - [45] G. Bakshi and N. Kapadia, Delta-hedged gains and the negative market volatility risk premium, *The Review of Financial Studies* **16**, 527 (2003).
 - [46] P. Carr and L. Wu, Variance risk premiums, *The Review of Financial Studies* **22**, 1311 (2009).
 - [47] O. Peters, The ergodicity problem in economics, *Nature Physics* **15**, 1216 (2019).
 - [48] T. M. Cover, Universal portfolios, *Mathematical Finance* **1**, 1 (1991).
 - [49] P. M. DeMarzo, I. Kremer, and Y. Mansour, Robust option pricing: Hannan and Blackwell meet Black and Scholes, *Journal of Economic Theory* **163**, 410 (2016).
 - [50] J. S. Bell, Against ‘measurement’, *Physics World* **3**, 33 (1990).
 - [51] J. L. Kelly, A new interpretation of information rate, *The Bell System Technical Journal* **35**, 917 (1956).
 - [52] R. M. Solow, Technical change and the aggregate production function, *The Review of Economics and Statistics* **39**, 312 (1957).
 - [53] J. A. Schumpeter, *Capitalism, Socialism and Democracy* (Harper and Brothers, New York, NY, 1942).
 - [54] S. J. Grossman and J. E. Stiglitz, On the impossibility of informationally efficient markets, *The American Economic Review* **70**, 393 (1980).
 - [55] P. Milgrom and N. Stokey, Information, trade and common knowledge, *Journal of Economic Theory* **26**, 17 (1982).
 - [56] A. Sandroni, Do markets favor agents able to make accurate predictions?, *Econometrica* **68**, 1303 (2000).
 - [57] A. W. Lo, The adaptive markets hypothesis: Market efficiency from an evolutionary perspective, *The Journal of Portfolio Management* **30**, 15 (2004).